

JXTA v2.0 Protocols Specification

Copyright © 2001, 2002, 2005, 2006, 2007 Sun Microsystems Inc.

Copyright © 2002, 2003, 2004 The Internet Society. All Rights Reserved.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this paragraph are included on all such copies and derivative works. However, this document itself may not be modified in any way, such as by removing the copyright notice or references to the Internet Society or other Internet organizations, except as needed for the purpose of developing Internet standards in which case the procedures for copyrights defined in the Internet Standards process must be followed, or as required to translate it into languages other than English.

The limited permissions granted above are perpetual and will not be revoked by the Internet Society or its successors or assigns.

This document and the information contained herein is provided on an "AS IS" basis and *THE INTERNET SOCIETY AND THE INTERNET ENGINEERING TASK FORCE DISCLAIMS ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.*

COLLABORATORS

	<i>TITLE :</i> JXTA v2.0 Protocols Specification		<i>REFERENCE :</i>
<i>ACTION</i>	<i>NAME</i>	<i>DATE</i>	<i>SIGNATURE</i>
WRITTEN BY		June 1, 2007	

REVISION HISTORY

NUMBER	DATE	DESCRIPTION	NAME
1.2	June 22, 2001	preliminary DocBook format	goober@jxta.org,
1.2.2	September 22, 2001	first release based on DocBook format (no intended changes from 1.2 revision)	goober@jxta.org
1.2.3	November 20, 2001	Updates to the IDs chapter to document new IDs design. Information on components provided as part of reference implementations is also being moved to the reference implementation section (this may eventually be a separate document).	bondolo@jxta.org
1.2.4	November 20, 2001	Updates to Protocol chapter, discovery behavior and minor changes to the Resolver section.	
1.2.5	November 21, 2001	Updates to Advertisements chapter, added sections for ModuleClassAdvertisement, ModuleSpecAdvertisement, ModuleImplAdvertisement. Updated sections for PeerAdvertisement and PeerGroupAdvertisement.	bondolo@jxta.org
1.2.6	November 28, 2001	Updates to Advertisements chapter, added sections for PeerInfoAdvertisement. also includes typo corrections in various sections	bondolo@jxta.org
1.2.7	November 29, 2001	Updates to the Protocol Section, and more description on PeerInfoAdvertisement traffic field	bondolo@jxta.org
1.2.8	December, 2001	Integrated docbookized version of Eric's sections about endpoint service and router. Made some cosmetic fixes per last review.	bondolo@jxta.org

REVISION HISTORY

NUMBER	DATE	DESCRIPTION	NAME
1.2.9	December 4, 2001	Updates to protocols chapter to use schemas to describe message structure more formally. Remove Peer Membership Protocol. Remove Content Advertisement (may be readed in documentation of CMS, but not part of or used by core protocols) Minor cosmetic improvements to copyright and ids chapter.	bondolo@jxta.org
1.2.10	December 7, 2001	Updates to protocols chapter to use mores schemas to describe message structure more formally. Examples for each as well.	bondolo@jxta.org
1.2.11	December 7, 2001	More updates to the protocols chapter. Converting raw program listings to figures and examples, addition of live cross references. Preparing for glossary and bibliography.	bondolo@jxta.org
1.2.12	December 9, 2001	More updates to the protocols chapter. PeerInfoProtocol and Pipe Binding Protocol now have schemas. Converting raw program listings to figures and examples, addition of live cross references. Preparing for glossary and bibliography. Cosmetic changes to Messages chapter	bondolo@jxta.org
1.2.13	December 17, 2001	Peer Info Protocol was redone because the previous schema though XML-like, was not representable with XML. (Endpoint Addresses as tag names were the culprit). PIP now looks more like the Peer Discovery Protocol with separate schemas for query and response. Document has been restructured into three major sections; Core Specification, Standard Services, and Reference Implementations. This restructuring makes understanding the JXTA Protocols more straight-forward because requirements and behaviours are now divided between functionality required by all JXTA peers, functionality needed to interoperate with the optional JXTA Protocols, and functionality required to interoperate with specific JXTA implementations.	bondolo@jxta.org
1.3	January 3, 2002	Wrote introductions for each of the new document sections. Introduced a glossary and bibliography. More restructuring of the source docbook files. Changes to prepare for a printable form document.	

REVISION HISTORY

NUMBER	DATE	DESCRIPTION	NAME
1.3.1	January 9, 2002	Spelling and grammar corrections. More reorganization of material into core, standard services and reference implementation parts. Updating stylesheets and docbook usage based on techniques seen in other docbook projects.	bondolo@jxta.org
1.3.2	January 9, 2002	More moving content around into core and standard services. More chapters reorganized. New text in advertisements chapter. More bibliographic information.	bondolo@jxta.org
1.3.3	January 14, 2002	Redid the schemas for the advertisement chapter. Moved all of the protocol advertisements and schemas to the appropriate protocols chapters. Use of <literal> has now been replaced with <phrase> and a role attribute of <i>rfc2119</i> . Improved field descriptions for some of the core advertisements.	bondolo@jxta.org
1.3.4	February 13, 2002	Lots of minor revisions based on review cycle. Thanks to all of the reviewers who took time to review and respond!	bondolo@jxta.org
1.3.5	April 19, 2002	Redid some of the schema descriptions in a new and clearer style, more will be converted over time. Redid schemas and examples for Pipe Binding Protocol section. A few changes to FO generation which generates <i>much</i> better PDF output. Corrected a few specific issues pointed out by Brendon Wilson.	bondolo@jxta.org
1.3.6	June 10, 2002	More complete documentation of the Pipe Binding Protocol based upon review of the J2SE and C implementations. Partial re-write of the Messages chapter to flesh out some of the details. Remove text not specific to the protocol (mostly dealing with namespace ":" name).	bondolo@jxta.org
1.3.7	June 11, 2002	Minor corrections following review. Now has an "Abstract" section.	bondolo@jxta.org
1.3.8	June 18, 2002	Lots of corrections following a detailed review by Jeff Altman, with many more corrections to come! A bit of reorganisation of the messages chapter into two pieces. Some cleanups for TOC generation in the HTML version. Still working on SVG figures. Next release will be 1.4 upon completion of all the feedback cleanups.	bondolo@jxta.org

REVISION HISTORY

NUMBER	DATE	DESCRIPTION	NAME
1.4	June 19, 2002	The Docbook source now validates against the Docbook 4.2 DTD. This seemingly minor issue fixed a number of problems with the HTML, PDF and rfc 2629 renderings and cross-referencing. Implemented all of the feedback from the most recent review cycle. Fixed a number of the figures.	bondolo@jxta.org
1.4.1	June 21, 2002	More revisions following careful review. This version is being released to coincide with the submission of the JXTA Specification as an Internet-Draft to IETF.	bondolo@jxta.org
1.4.2	October 10, 2002	Updates to Standard Transports sections. Updates to the Endpoint Routing Protocol and Endpoint Router Transport. Adds new optional SRDI message for Endpoint Resolver Protocol. Minor textual improvements and clarifications elsewhere. New examples and figures.	bondolo@jxta.org
2.0	February 27, 2003	Updates to Standard Transports sections. Clarifications in the IDs section. Updates to the Pipe Resolver protocol. Minor textual improvements and clarifications elsewhere. Added a section on "Users and Peers". Added a number of glossary terms. More contributions to the glossary are welcome. Updated version to "2.0". Rewrote or revised much of the Overview chapter.	bondolo@jxta.org
2.1.1	October 15, 2003	Additional updates to the Introduction and Overview chapters. Many additions to the glossary section. Minor schema corrections.	bondolo@jxta.org
2.2.1	April 10th, 2004	Additional updates to glossary and bibliography. Added updates to pipes regarding propagate pipes.	bondolo@jxta.org
2.3	June 30th, 2004	Additional information about the Rendezvous, Resolver and Pipe Services.	bondolo@jxta.org
2.3.5	September 20th, 2005	Homogeneous spec for Resolver queries and responses.	mathieu@jxta.org
2.5	June 23rd, 2006	Clarifications to the HTTP Message Transport.	bondolo@jxta.org

REVISION HISTORY

NUMBER	DATE	DESCRIPTION	NAME
2.5.1	August 23rd, 2006	Corrections and clarifications to Route Advertisement and Pipe Resolver Message.	bondolo@jxta.org
2.5.2	January 9th, 2007	Corrections and clarifications from several sharp eyed readers. Now includes complete specification of Binary Message Format version 2 and specification of the endpoint SrcAddress and DstAddress message elements.	bondolo@jxta.org

Contents

I	JXTA Core Specification	14
1	IDs	15
1.1	Introduction	15
1.2	JXTA ID Properties	15
1.3	Using JXTA IDs in Protocols	15
1.4	Format of a JXTA ID URN	16
1.5	Example JXTA ID URNs	16
1.6	JXTA ID Representation	16
1.7	JXTA ID Formats	17
1.8	JXTA ID Types	17
1.8.1	Peer Group IDs	18
1.8.2	Peer IDs	18
1.8.3	Codat IDs	18
1.8.4	Pipe IDs	18
1.8.5	Module Class IDs	18
1.8.6	Module Spec IDs	18
1.9	JXTA ID Formats : 'jxta' ID Format	18
2	Advertisements	20
2.1	Introduction	20
2.2	XML and JXTA Advertisements	20
2.3	Peer Advertisement	21
2.4	Peer Group Advertisement	22
2.5	Module Class Advertisement	23
2.6	Module Specification Advertisement	23
2.7	Module Implementation Advertisement	25

3	JXTA Core Protocols	27
3.1	Peer Resolver Protocol	27
3.1.1	Introduction	27
3.1.2	Resolver Query Message	27
3.1.3	Resolver Response Message	29
3.1.4	Resolver SRDI Message	32
3.1.5	Listener and Element Naming	32
3.1.6	Behaviour	33
3.1.6.1	Handler Name	33
3.1.6.2	Policies and Quality of Service	34
3.2	Endpoint Routing Protocol	34
3.2.1	Endpoint Addresses	35
3.2.2	Route information	36
3.2.3	Route Query Message	37
3.2.4	Route Response Message	37
4	Core JXTA Message Transport Bindings	39
4.1	Endpoint Service	39
4.1.1	Description	39
4.1.2	Protocol	39
4.2	Endpoint Router Transport Protocol	40
4.2.1	Description	40
4.2.2	Protocol	40
4.2.3	EndpointRouter Message Element	40
4.2.3.1	EndpointRouter Endpoint Address format	41
5	Messages	42
5.1	Introduction	42
5.2	Message	42
5.3	Element	42
5.3.1	Namespace	42
5.3.2	Name	43
5.3.3	Type	43
5.3.4	Content	43

II	JXTA Standard Services	44
6	Standard Protocols	45
6.1	Peer Discovery Protocol	45
6.1.1	Introduction	45
6.1.2	Discovery Query Message	45
6.1.3	Discovery Response Message	47
6.1.4	Behaviour	49
6.2	Rendezvous Protocol	50
6.2.1	Introduction	50
6.2.2	Rendezvous Advertisement	50
6.2.3	PeerView Protocol	51
6.2.4	Rendezvous Lease Protocol	53
6.2.4.1	Lease Request Message	53
6.2.4.2	Lease Granted Message	53
6.2.4.3	Lease Cancel Message	53
6.2.5	Message Propagation Protocol	53
6.2.6	Behaviour	53
6.2.6.1	Peer connection	53
6.2.6.2	Propagation Control	54
6.3	Peer Information Protocol	54
6.3.1	Obtaining PIP Responses	54
6.3.2	PIP Query Message	55
6.3.3	PIP Response Message	55
6.4	Pipe Binding Protocol	57
6.4.1	Pipe Advertisement	57
6.4.2	Pipe Resolver Message	58
6.4.3	Propagate Pipe Message Header	60
7	Standard JXTA Message Transports	61
7.1	TCP/IP Message Transport	61
7.1.1	Introduction	61
7.1.2	Choreography	61
7.1.3	Welcome Message	61
7.1.4	Message Transmission	62
7.1.4.1	JXTA Message Package - Required Headers	63
7.1.4.2	JXTA Message Package - Optional Headers	63
7.1.5	IP Multicast Usage	63
7.2	HTTP Message Transport	64

7.2.1	The HTTP Initiator	64
7.2.2	The HTTP Receiver	64
7.2.3	HTTP Messages	64
7.2.3.1	‘Ping’ Command	64
7.2.3.2	‘Poll’ Command	65
7.2.3.3	‘Send’ Command	66
7.3	TLS Transport Binding	67
7.3.1	Introduction	67
7.3.2	TLS Messages	67
7.3.2.1	TLS Content Element	67
7.3.2.2	TLS ACK Element	67
7.3.2.3	Retry Element	68
8	JXTA Message Wire Representations	69
8.1	General Requirements	69
8.2	Binary Message Format	69
8.2.1	Conventions	69
8.2.2	Binary Message Version 1	70
8.2.2.1	Message Header	71
8.2.2.2	Namespaces	71
8.2.2.3	Message Element Header	71
8.2.3	Binary Message Version 2	71
8.2.3.1	Message Header	73
8.2.3.2	Names Table	73
8.2.3.3	Message Element Header	73
8.3	XML Message Format	73
8.3.1	Message	74
8.3.2	Element	74
8.3.2.1	Name	74
8.3.2.2	MIME type	74
8.3.2.3	Encoding	74
III	JXTA Reference Implementations Information	75
9	Java 2 SE Reference Implementation	76
9.1	JXTA ID Formats	76
9.1.1	JXTA ID Formats : ‘uuid’ ID Format	76
9.1.2	JXTA UUID Field Definitions	77
9.1.2.1	JXTA UUID Codat ID Fields	77

9.1.2.2	JXTA UUID PeerGroup ID Fields	77
9.1.2.3	JXTA UUID Peer ID Fields	78
9.1.2.4	JXTA UUID Pipe ID Fields	78
9.1.2.5	JXTA UUID Module Class ID Fields	79
9.1.2.6	JXTA UUID Module Spec ID Fields	79
9.2	J2SE JXTA Endpoint Router Implementation	79
10	Glossary	81
11	Bibliography	83

List of Figures

1	JXTA Protocols	4
2	JXTA Pipe Types	12
1.1	JXTA ID ABNF	17
1.2	JXTA ID : "jxta" ID Format ABNF	19
2.1	Common Advertisement Fragments Schemas	21
2.2	Peer Advertisement Schema	21
2.3	Peer Group Advertisement Schema	22
2.4	Module Class Advertisement Schema	23
2.5	Module Specification Advertisement Schema	24
2.6	Module Implementation Advertisement Schema	25
3.1	Resolver Query Schema	28
3.2	Resolver Response Schema	30
3.3	Resolver SRDI Schema	32
3.4	Listener Naming Syntax ABNF	33
3.5	Endpoint Address URI ABNF	35
3.6	Route Advertisement	36
3.7	Endpoint Router Query	37
3.8	Endpoint Router Response Message	37
4.1	Endpoint Router Message Element	41
4.2	JXTA Endpoint Router Address Format	41
6.1	Discovery Query Schema	46
6.2	Discovery Response Schema	48
6.3	Rendezvous Advertisement Schema	50
6.4	PeerView Message	51
6.5	PeerView Message Scenarios	52
6.6	RendezVous Propagate Message Schema	54
6.7	PIP Query Message	55

6.8	PIP Response Message	56
6.9	Pipe Advertisement Schema	58
6.10	Pipe Resolver Message Schema	59
6.11	Propagate Pipe Message Header Schema	60
7.1	Welcome Message ABNF	62
7.2	JXTA Message Package ABNF	63
8.1	JXTA Binary Message Version 1 ABNF	70
8.2	JXTA Binary Message Version 2 ABNF	72
9.1	JXTA "uuid" ID Format ABNF	76
9.2	JXTA UUID Codat ID Fields	77
9.3	JXTA UUID PeerGroup ID Fields	78
9.4	JXTA UUID Peer ID Fields	78
9.5	JXTA UUID Pipe ID Fields	78
9.6	JXTA UUID Module Class ID Fields	79
9.7	JXTA UUID Module Spec ID Fields	79

List of Examples

1.1	Sample JXTA IDs	16
3.1	Resolver Query	29
3.2	Resolver Response	31
3.3	Resolver SRDI	32
3.4	Endpoint Address URI Examples	35
6.1	Discovery Query	47
6.2	Discovery Response	49
6.3	Pipe Advertisement	58
7.1	Welcome Message	62
7.2	HTTP Ping Example	65
7.3	HTTP Poll Example	66
7.4	HTTP Send Example	67
8.1	XML Message	74
9.1	Sample "uuid" Format ID	77

Abstract

The JXTA protocols defines a suite of six XML-based protocols that standardize the manner in which peers self-organize into peer groups, publish and discover peer resources, communicate, and monitor each other. The Endpoint Routing Protocol (ERP) is the protocol by which a peer can discover a route (sequence of hops) to send a message to another peer potentially traversing firewalls and NATs. The Rendezvous Protocol (RVP) is used for propagating a message within a peer group. The Peer Resolver Protocol (PRP) is the protocol used to send a generic query to one or more peers, and receive a response (or multiple responses) to the query. The Peer Discovery Protocol (PDP) is used to publish and discover resource advertisements. The Peer Information Protocol (PIP) is the protocol by which a peer may obtain status information about another peers. The Pipe Binding Protocol (PBP) is the protocol by which a peer can establish a virtual communication channel or pipe between one or more peers. The JXTA protocols permit the establishment a virtual network overlay on top of physical networks allowing peers to directly interact and organize independently of their network location and connectivity. The JXTA protocols have been designed to be easily implemented on unidirectional links and asymmetric transports.

Conventions

The following conventions are used throughout this document.

Significant Keywords

The key words must, must not, required, shall, shall not, should, should not, recommended, not recommended, may, and optional in this document are to be interpreted as described in *IETF RFC 2119* "[RFC2119](#)".

Text Representations

All strings and text found in the JXTA Protocols should be assumed to be encoded using Unicode "[USA28](#)" Canonical UTF8 (NFC) (see *Unicode Standard Annex #15 : Unicode Normalization Forms* "[USA15](#)") unless otherwise specified.

Introduction

Why JXTA?

The [JXTA](#) Protocols comprise an open network computing platform designed for peer-to-peer (P2P) computing. The set of generalized JXTA protocols enable all connected devices on the network -- including cell phones, PDAs, PCs and servers -- to communicate and collaborate as peers. The JXTA protocols enable developers to build and deploy interoperable services and applications, further spring-boarding the peer-to-peer revolution on the Internet.

The JXTA protocols standardize the manner in which peers:

- Discover each other
- Self-organize into peer groups
- Advertise and discover network resources
- Communicate with each other
- Monitor each other

The JXTA protocols are designed to be independent of the underlying implementation. In particular, the JXTA protocols *do not*:

- Require the use of any particular computer language or operating system.
- Require the use of any particular network transport or topology.
- Require the use of any particular authentication, security or encryption model.

JXTA provides a simple and generic P2P platform with all the basic functions necessary to host all types of network services:

- JXTA is defined by a small number of protocols. Each protocol is easy to implement and integrate into P2P services and applications. Thus, service offerings from one vendor can be used transparently by the user community of another vendor's system.
 - The JXTA protocols are defined to be independent of programming languages, so that they can be implemented in C/C++, Java, Perl, and numerous other languages. Heterogeneous devices with completely different software stacks can interoperate using the JXTA protocols.
 - The JXTA protocols are designed to be independent of transport protocols. They can be implemented on top of TCP/IP, HTTP, Bluetooth, HomePNA, and many other protocols.
-

The JXTA Protocols

The JXTA protocols are a set of six protocols that have been specifically designed for ad hoc, pervasive, and multi-hop peer-to-peer (P2P) network computing. Using the JXTA protocols, peers can cooperate to form self-organized and self-configured peer groups independent of their positions in the network (edges, firewalls, network address translators, public vs. private address spaces), and without the need of a centralized management infrastructure.

The JXTA protocols are designed to have very low overhead, make few assumptions about the underlying network transport and impose few requirements on the peer environment, and yet are able to be used to deploy a wide variety of P2P applications and services in a highly unreliable and changing network environment.

Peers use the JXTA protocols to advertise their resources and to discover network resources (services, pipes, etc.) available from other peers. Peers form and join peer groups to create special relationships. Peers cooperate to route messages allowing for full peer connectivity. The JXTA protocols allow peers to communicate without the need to understand or manage the potentially complex and dynamic network topologies which are increasingly common.

The JXTA protocols allow peers to dynamically route messages across multiple network hops to any destination in the network (potentially traversing firewalls). Each message carries with it either a complete or partially ordered list of gateway peers through which the message might be routed. Intermediate peers in the route may assist the routing by using routes they know of to shorten or optimize the route a message is set to follow.

The JXTA protocols are composed of six protocols that work together to allow the discovery, organization, monitoring and communication between peers:

- **Peer Resolver Protocol (PRP)** is the mechanism by which a peer can send a query to one or more peers, and receive a response (or multiple responses) to the query. The PRP implements a query/response protocol. The response message is matched to the query via a unique id included in the message body. Queries can be directed to the whole group or to specific peers within the group.
- **Peer Discovery Protocol (PDP)** is the mechanism by which a peer can advertise its own resources, and discover the resources from other peers (peer groups, services, pipes and additional peers). Every peer resource is described and published using an advertisement. Advertisements are programming language-neutral metadata structures that describe network resources. Advertisements are represented as XML documents.
- **Peer Information Protocol (PIP)** is the mechanism by which a peer may obtain status information about other peers. This can include state, uptime, traffic load, capabilities, and other information.
- **Pipe Binding Protocol (PBP)** is the mechanism by which a peer can establish a virtual communication channel or pipe between one or more peers. The PBP is used by a peer to bind two or more ends of the connection (pipe endpoints). Pipes provide the foundation communication mechanism between peers.
- **Endpoint Routing Protocol (ERP)** is the mechanism by which a peer can discover a route (sequence of hops) used to send a message to another peer. If a peer 'A' wants to send a message to peer 'C', and there is no known direct route between 'A' and 'C', then peer 'A' needs to find intermediary peer(s) who will route the message to 'C'. ERP is used to determine the route information. If the network topology changes and makes a previously used route unavailable, peers can use ERP to find an alternate route.
- **Rendezvous Protocol (RVP)** is the mechanism by which peers can subscribe or be a subscriber to a propagation service. Within a peer group, peers can be either rendezvous peers or peers that are listening to rendezvous peers. The Rendezvous Protocol allows a peer to send messages to all the listening instances of the service. The RVP is used by the Peer Resolver Protocol and by the Pipe Binding Protocol in order to propagate messages.

All of these protocols are implemented using a common messaging layer. This messaging layer is what binds the JXTA protocols to various network transports. (see [Messages](#))

Each of the JXTA protocols is independent of the others. A peer is not required to implement all of the JXTA protocols to be a JXTA peer. A peer only implements the protocols that it needs to use. For example:

- A device may have all the necessary advertisements it uses pre-stored in memory, and therefore not need to implement the Peer Discovery Protocol.

- A peer may use a pre-configured set of router peers to route all its messages. Because the peer just sends messages to the known routers to be forwarded, it does not need to fully implement the Endpoint Routing Protocol.
- A peer may not need to obtain or wish to provide status information to other peers, hence the peer might not implement the Peer Information Protocol.

Each peer must implement two protocols in order to be addressable as a peer: the Peer Resolver Protocol and the Endpoint Routing Protocol. These two protocols and the advertisements, services and definitions they depend upon are known as the *JXTA Core Specification*. The *JXTA Core Specification* establishes the base infrastructure used by other services and applications.

The remaining JXTA protocols, services and advertisements are optional. JXTA implementations are not required to provide these services, but are strongly recommended to do so. Implementing these services provides greater interoperability with other implementations and broader functionality. These common JXTA services are known as the *JXTA Standard Services*.

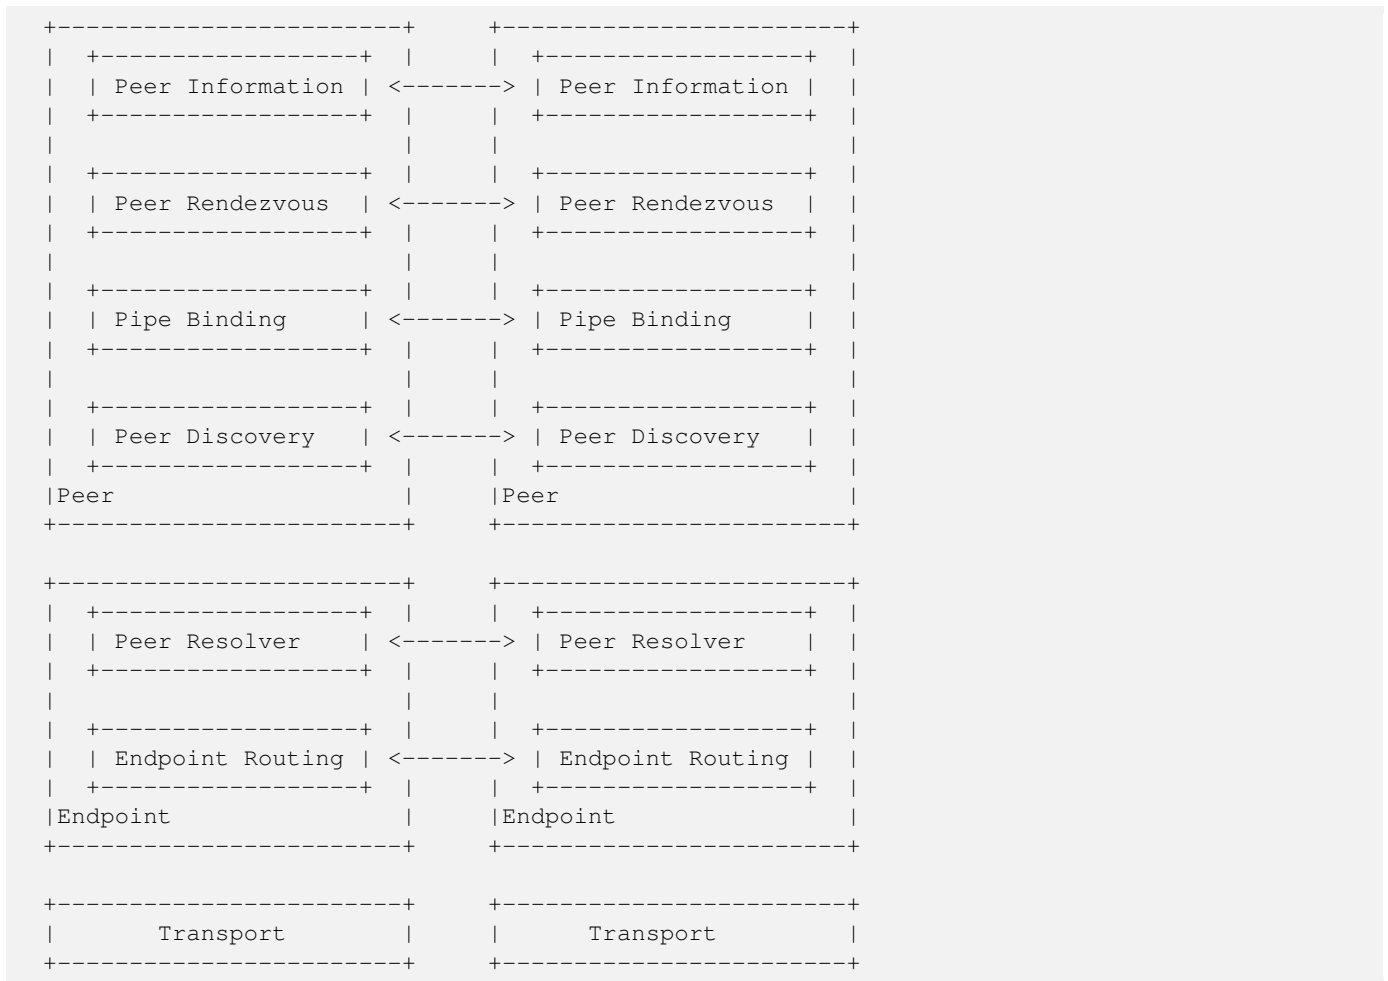


Figure 1: JXTA Protocols

A peer may decide to cache advertisements discovered via the Peer Discovery Protocol for later usage. It is important to point out that caching is not required by the JXTA architecture, but caching can be an important optimization. By caching advertisements, a peer avoids the need to perform a new discovery each time it accesses a network resource. In highly-transient environment, performing the discovery is the only viable solution. In static environments, caching is more efficient.

A unique characteristic of P2P networks, like JXTA, is their ability to spontaneously replicate information toward end-users. Popular advertisements are likely to be replicated more often, making them easier to find as more copies become available. Peers do not have to return to the same peer to obtain the advertisements they seek. Instead of querying the original source of an

advertisement, peers may query neighboring peers that have already cached the information. Each peer may potentially become an advertisement provider to any other peer.

The JXTA protocols have been designed to allow JXTA to be easily implemented on uni-directional links and asymmetric transports. In particular, many forms of wireless networking do not provide equal capability for devices to send and receive. JXTA permits any uni-directional link to be used when necessary, improving overall performance and network connectivity in the system. The intent is for the JXTA protocols to be as pervasive as possible, and easy to implement on any transport. Implementations on reliable and bi-directional transports such as TCP/IP or HTTP should lead to efficient bi-directional communications.

The JXTA uni-directional and asymmetric transport also plays well in multi-hop network environments where the message latency may be difficult to predict. Furthermore, peers in a P2P network tend to have nondeterministic behaviours. They may join or leave the network on a very frequent basis.

The JXTA Three Layer Cake

The JXTA Project is logically divided in three layers.

- *Platform.* This layer encapsulates minimal and essential primitives that are common to P2P networking, including peers, peer groups, discovery, communication, monitoring, and associated security primitives. This layer is ideally shared by all P2P devices so that interoperability becomes possible.
- *Services.* This layer includes network services that may not be absolutely necessary for a P2P network to operate but are common or desirable for P2P environments. Examples of network services include search and indexing, directory, storage systems, file sharing, distributed file systems, resource aggregation and renting, protocol translation, authentication and PKI services.
- *Applications.* This layer includes P2P instant messaging, entertainment content management and delivery, P2P E-mail systems, distributed auction systems, and many others. Obviously, the boundary between services and applications is not rigid. An application to one customer can be viewed as a service to another customer.

JXTA Assumptions

This section is a guide to the assumptions that shape the JXTA design. There are two types of assumptions stated here: those which describe the requirements of JXTA implementations and those which describe the expected behaviour of the JXTA network.

A peer shall not make assumptions about the runtime environments or programming languages in use by another peer. The network of peers reachable by any peer is likely to contain many peers with heterogeneous implementations and capabilities.

A peer should assume that the capabilities and complexity of the network peers supporting these protocols can range from a single light switch to a highly-available supercomputer cluster.

A peer must implement the JXTA protocols such that all interaction with other peers is correct and conformant.

A peer may implement only the JXTA protocols it requires for correct and conformant interaction with other peers.

A peer may choose to partially implement protocols if unimplemented portions will never be used. (e.g. it may implement client-side or server-side portions only.)

Peers wishing to interact with other peers within the JXTA network should be willing to participate fully in the protocols. In particular, peers should cache advertisements and forward messages for other peers in the JXTA network. However, this participation is optional.

The JXTA protocols may be deployed over a wide variety of network configurations including the Internet, corporate intranets, a dynamic proximity network, or in a home networking environment. Applications should avoid assumptions about the underlying network environment.

Peers receiving a corrupted or detectably compromised message must discard the message. Messages may be corrupted or intentionally altered during network transmission.

Peers may appear, disappear and migrate at any time without notice. In particular, the JXTA protocols support very arbitrary environment changes allowing a peer to dynamically discover and reconnect to its changing environment.

The communication path between any pair of peers may at times not work equally well in both directions. That is, communications between two peers will in many cases be able to operate bi-directionally, but at times the connection between two peers may be only uni-directional, allowing one peer to successfully send messages to the other while no communication is possible in the reverse direction.

The JXTA protocols are defined as idempotent protocol exchanges. The same messages may be sent/received more than once during the course of a protocol exchange. No protocol states are required to be maintained at both ends.

Due to unpredictability of P2P networks, assumptions must not be made about the time required for a message to reach a destination peer. The [JXTA Core Protocols](#) must not impose any timing requirements for message receipt.

The JXTA Transport Protocols (see [Core JXTA Message Transport Bindings](#) and [Standard JXTA Message Transports](#)) must not impose any timing requirements on the [JXTA Core Protocols](#), but may have timing requirements internal to themselves.

The [Standard Protocols](#) (e.g. [Peer Discovery Protocol](#), [Peer Information Protocol](#), [Peer Discovery Protocol](#), etc.) and application defined protocols may impose timing requirements on message delivery and receipt.

A JXTA protocol message which is routed through multiple hops should not be assumed to reliably delivered, even if only reliable transports such as TCP/IP are used through all hops. A congested peer may drop messages at any time rather than routing them.

JXTA protocol messages and advertisements are defined using XML and are required to be well-formed XML documents. The use of a full XML parser is optional as long as the requirements of the being well-formed and the JXTA protocol are met. Small JXTA implementations may choose to use pre-built XML or XML templates for message and advertisement construction.

The JXTA protocols must not require a broadcast or multicast capability of the underlying network transport. Messages intended for receipt by multiple peers (propagation) may be implemented using point-to-point communications.

A peer should make the assumption that if a destination address is not available at any time during the message transmission, the message will not be delivered.

A peer must not assume that there is a guaranteed return route to a peer from which it has received communication. The lack of a return route may either be temporary or permanent.

Each peer must be a member of the World Peergroup and Net Peergroups. Membership in these groups is automatic.

Peers must be members of the same peer group in order to exchange messages.

Names are not unique unless a coordinated naming service is used to guarantee name uniqueness. A naming service is typically a service that guarantees, within a given scope, the uniqueness of name and can be used to register name mapping. Examples of name services are DNS and LDAP. A naming service is optional. JXTA does not define its own naming service.

Once content has been published to the JXTA network, it should not be assumed that that the content can be later retrieved from the JXTA network. The content may be available only from peers that are not currently reachable.

Once a content has been published to the JXTA network, it must not be assumed that the content can be deleted. Republication of content by peers is unrestricted and the content may propagate to peers which are not reachable from the publishing peer.

Conceptual Overview

JXTA is intended to be a small system with a limited number of concepts at its core. This chapter introduces the concepts which are central to JXTA.

Peers

A JXTA peer is any networked device (sensor, phone, PDA, PC, server, supercomputer, etc.) that implements the core JXTA protocols. Each peer is identified by a unique ID. Peers are autonomous and operate independently and asynchronously of all other peers. Some peers may have dependencies upon other peers due to special requirements such as the need for gateways, proxies, or routers.

Peers may publish network services and resources (CPU, storage, databases, documents, etc.) for use by other peers. A peer may cache advertisements for JXTA resources, but doing so is optional. Peers may have persistent storage.

Peers are typically configured to spontaneously discover each other on the network to form transient or persistent relationships with other peers. Peers that provide the same set of services tend to be interchangeable. As a result, peers typically need to interact with only a small number of other peers (network neighbors or buddy peers). Peers should not make assumptions about the reliability of other peers. Peers may join or leave the network at any time. A peer should always anticipate that connectivity may be lost to any peer that it is currently communicating with.

Peers may advertise multiple network interfaces. Each published interface is advertised as a peer endpoint. A peer endpoint is a URI that uniquely identify a peer network interface (for example, a URI might specify a TCP port and associated IP address). Peer endpoints are used by peers to establish direct point-to-point connections between two peers.

Communicating peers are not required to have direct point-to-point network connection between themselves. A peer may need to use one or more intermediary peers to route a message to another peer that is separated due to physical network connections or network configurations (e.g., NATs, firewalls, or proxies).

Peers and Users

There is no explicit relationship between users and peers within JXTA, but for many applications the two concepts are closely related.

Frequently a peer will be under the control of a human operator, the ‘user’, and will interact with the network on the basis of the user’s direction. However, not every peer has an associated user. Many peers exist to provide services and resources to the network, but are not associated with any user. Examples include devices such as sensors and printers, and services such as databases.

The concept of ‘user’ exists primarily for identification: some applications may need to allow users to tell each other apart and to manage their interactions. Identity is also frequently tied to security and permissions.

Peer Groups

Peers self-organize into Peer Groups. A peer group is a collection of peers that have a common set of interests. Each peer group is uniquely identified by a PeerGroup Id. The JXTA protocols do not dictate when, where, or why peer groups are created. The JXTA protocols only describe how peers may publish, discover, join, and monitor peer groups.

JXTA recognizes three common motivations for creating peer groups:

- *To create a secure environment.* Peer group boundaries permit member peers to access and publish protected contents. Peer groups form logical regions whose boundaries limit access to the peer group resources. A peer group does not necessarily reflect the underlying physical network boundaries such as those imposed by routers and firewalls. Peer groups virtualize the notion of routers and firewalls, subdividing the network into secure regions without respect to actual physical network boundaries.
- *To create a scoping environment.* Peer groups are typically formed and self-organized based upon the mutual interest of peers. No particular rules are imposed on the way peer groups are formed, but peers with the same interests will tend to join the same peer groups. Peer groups serve to subdivide the network into abstract regions providing an implicit scoping mechanism. Peer group boundaries define the search scope when searching for a group's content.
- *To create a monitoring environment.* Peer groups permit peers to monitor a set of peers for any special purpose, including heartbeat, traffic introspection, and accountability.

A peer group provides a set of services called peer group services. JXTA defines a core set of peer group services, and the JXTA protocols specify the wire format for these services. Additional peer group services can be developed for delivering specific services. For example, a lookup service could be implemented to find active (running on some peer) and inactive (not yet running) service instances. The core peer group services are:

Discovery Service

The Discovery service is used by member peers to search for peer group resources such as peers, peer groups, pipes, and services.

Membership Service

The membership service is used by the current peer group members to reject or accept a new group membership application (i.e., to allow a new peer to join a peer group). We expect that most peer groups will have a membership service, though it may be a 'null' authenticator service which imposes no real membership policy. Peers wanting to join a peer group may need to discover at least one member of the peer group and then request to join. The request to join is either rejected or accepted by the collective set of current members. The membership service may enforce a vote of peers or elect a designated group representative to accept or reject new membership requests. A peer may belong to more than one peer group simultaneously.

Access Service

The Access service is used to validate requests made by one peer to another. The peer receiving the request provides the requesting peer credentials and information about the request being made to the Access Service to determine if the access is permitted. Not all actions within the peer group need to be checked with the Access Service. Only those actions that are restricted to a subset of member peers must be checked.

Pipe Service

The pipe service is used to manage and create pipe connections between the different peer group members.

Resolver Service

The resolver service is used to address queries to services running on peers in the group and collect responses.

Monitoring Service

The monitoring service is used to allow one peer to monitor other members of the same peer group.

Not all the above services must be implemented by a peer group. A peer group can implement only the services it finds useful and rely on the default Net Peergroup to provide generic implementations of other services. Each service may implement one or more of the JXTA protocols, the specifications for which are the main content of this document. A service will typically implement one protocol for simplicity and modularity reasons, but some services may not implement any protocols.

Network Services

Peers cooperate and communicate to publish, discover and invoke network services. A peer can publish as many services as it can provide. Peers discover network services via the Peer Discovery Protocol.

Network Service specifications are beyond the scope of this document. Upcoming standards such as WSDL, ebXML, SOAP, UPnP might be used within a JXTA Network.

The JXTA protocols recognize two levels of network services:

- Peer Services
- Peer Group Services

A peer service is accessible only on the peer that is publishing the service. If that peer happens to fail, then the service also fails. Multiple instances of the service can be run on different peers, but each instance publishes its own advertisement.

A peer group service is composed of a collection of instances (potentially cooperating with each other) of the service running on multiple members of the peer group. If any one peer fails, the collective peer group service is not affected: in most cases, the service is still available from another peer member. Peer group services are published as part of the peer group advertisement.

Services can either be pre-installed into a peer or loaded from the network. The process of finding, downloading and installing a service from the network is similar to performing a search on the internet for a web page, retrieving the page, and then installing the required plug-in. In order to actually run a service, a peer may need to locate an implementation suitable for the peer's runtime environment. Multiple implementations of the same service may allow Java peers to use Java code implementations, and native peers to use native code implementations.

Service Invocation

Service invocation is outside the scope of JXTA. JXTA is designed to interoperate and be compatible with any Web service standards: WSDL, uPnP, RMI, etc. The JXTA protocols define a generic framework to publish and discover advertisements that may describe services. Peers publish and discover advertisements via the Peer Discovery Protocol. An advertisement for a service typically contains all necessary information to either invoke or instantiate the service being described. The JXTA protocols define Module Advertisements, but any other form of service description may be introduced.

JXTA-Enabled Services

JXTA-Enabled services are services that are published using a ModuleSpecAdvertisement. A ModuleSpecAdvertisement may specify a pipe advertisement that can be used by a peer to create output pipes to invoke the service. ModuleSpecAdvertisements may be extended in the future to contain a list of pre-determined messages that can be sent by a peer to interact with the service. A ModuleSpecAdvertisement may also contain references to two other services that can be used as an authenticator for the service and as a local proxy for the service.

Each JXTA-enabled service is uniquely identified by its ModuleSpecID.

IDs

Peers, peer groups, pipes, contents, and other resources need to be uniquely identifiable within the JXTA protocols. A JXTA ID uniquely identifies an entity and serves as a canonical means of referring to that entity.

URNs are used for the expression of JXTA IDs.

Advertisements

All network resources, such as peers, peer groups, pipes and services, are represented by advertisements. Advertisements are JXTA's language neutral metadata structures for describing resources. The JXTA Core Specification and Standard Services define, among others, the following advertisement types:

- Peer Advertisement

- PeerGroup Advertisement
- ModuleClass Advertisement
- ModuleSpec Advertisement
- ModuleImpl Advertisement
- Pipe Advertisement
- Rendezvous Advertisement

The complete specification of advertisements is given in the [Advertisements](#) chapter. Because the JXTA protocols make heavy reference to these advertisements, the reader should be familiar with advertisements before moving on to the protocol specification chapters. Advertisements are, by far, the most common document exchanged in the protocol messages.

Advertisements are published with a lifetime that specifies the availability of the resource. An advertisement can be republished to extend the lifetime of a resource.

Services and peer implementations can create their own advertisement types, either from scratch or by subtyping the existing types.

Credentials

The need to support different levels of resource access in a dynamic and ad hoc P2P network leads to a role-based trust model in which an individual peer acts under the authority granted to it by another trusted peer to perform a particular task. Peer relationships may change quickly and the policies governing access control need to be flexible in allowing or denying access.

Four basic security requirements must be provided:

Confidentiality

Guarantees that the contents of the message are not disclosed to unauthorized individuals.

Authorization

Guarantees that the sender is authorized to send a message.

Data integrity

Guarantees that the message was not accidentally or deliberately modified in transit.

Refutability

Guarantees the message was transmitted by a properly identified sender and is not a replay of a previously transmitted message.

The structure of JXTA messages enables JXTA applications to add arbitrary metadata information, such as credentials, digests, certificates, and public keys, to messages. Message digests guarantee the data integrity of messages. Messages may also be encrypted and signed for confidentiality and refutability. Credentials can be used to provide message authentication and authorization.

A credential is a token that, when presented in a message body, is used to identify a sender and can be used to verify that sender's right to send the message to the specified endpoint. The credential is an opaque token that must be presented each time a message is sent. The sending address placed in the message envelope is cross-checked with the sender's identity in the credential. Each credential's implementation is specified as a plug-in configuration, which allows multiple authentication configurations to co-exist on the same network.

It is the intent of the JXTA protocols to be compatible with widely accepted transport-layer security mechanisms. Some JXTA implementations contain a virtualized TLS implementation that allows it to secure endpoint-to-endpoint communications regardless of the number of hops required to deliver each message.

TLS and IPSec could also be used as JXTA transports. However, when used as transports they provide integrity and confidentiality of message transfer only between the two communicating peers.

Pipes

Pipes are virtual communication channels used to send and receive messages between services or applications. Pipes provide the illusion of a virtual in and out mailbox that is independent of any single peer location, and independent of network topology (multi-hops route).

Different quality of services can be implemented by a pipe. For example:

Uni-directional asynchronous

The endpoint sends a message, no guarantee of delivery is made.

Synchronous request-response

The endpoint sends a message, and receives a correlated answer.

Bulk transfer

Bulk reliable data transfer of binary data.

Streaming

Efficient control-flow data transfer.

Secure

Secure reliable data streams.

The uni-directional asynchronous pipe is required by the JXTA protocols. Other pipe variations may be implemented for use by services and their associated protocols.

Pipes provide a network abstraction over the peer endpoint transport. Peer endpoints correspond to the available peer network interfaces that can be used to send and receive data from another peer. Pipes connect one or more peer endpoints. At each endpoint software to send or receive, as well as to manage associated pipe message queues or streams, is assumed, but message queues are optional.

The pipe endpoints are referred to as input pipes (receiver) and output pipes (transmitter). Pipe endpoints are dynamically bound to a peer endpoint at runtime, via the [Pipe Binding Protocol](#). The pipe binding process consists of searching for and connecting the two or more endpoints of a pipe.

When a message is sent into an output pipe, the message is sent by the local peer endpoint to the peer endpoint(s) where the associated input pipe is located. The set of peer endpoints currently associated with the input pipes is discovered using the [Pipe Binding Protocol](#).

A pipe offers two modes of communication:

Point to Point

A point to point pipe connects exactly two pipe endpoints together: an input pipe that receives messages sent from an output pipe. No reply or acknowledgment operation is supported. Additional information in the message payload (like a unique id) may be required to thread message sequences. The message payload may also contain a pipe advertisement that can be used to open a pipe to reply to the sender (send/response).

Propagate Pipe

A propagate pipe connects one output pipe and multiple input pipes together. Messages flow from the output pipe (propagation source) into the input pipes. A propagate message is sent to all listening input pipes. This process may create multiple copies of the message to be sent. On TCP/IP, when the propagate scope maps an underlying physical subnet in a 1-to-1 fashion, IP multicast may be used as an implementation for propagate pipes. For transports (such as HTTP) that do not provide multicast, propagate pipes can be implemented using point-to-point communication.




Figure 2: JXTA Pipe Types

Pipes may connect two peers that do not have a direct physical link. In this case, one or more intermediary peer endpoints are used to route messages between the two pipe endpoints.

Messages

The information transmitted using pipes and between endpoints is packaged as messages. Messages define an envelope to transfer any kind of data. A message may contain an arbitrary number of named sub-sections which can hold any form of data.

It is the intent that the JXTA protocols be compliant with W3C XML Protocol standards, so that the JXTA protocols can be implemented on XML transports such as SOAP, XML-RPC, etc.

The JXTA protocols are specified as a set of XML messages exchanged between peers. Each software platform binding describes how a message is converted to and from a native data structures such as Java objects or 'C' structures.

The use of XML messages to define protocols allows many different kinds of peers to participate in a protocol. Each peer is free to implement the protocol in a manner best suited to its abilities and role.

Codats

The JXTA protocols assume that many types of contents may be shared, exchanged, and replicated between peers. A content can be a text file, a structured document (like a PDF or an XML file), a Java `.jar` or loadable library, code or even an executable process. No size limitation is assumed.

Content are published and shared amongst peer members of a peer group. When a content is shared within the JXTA network it is associated with a JXTA ID. The combination of a content and a JXTA ID is known as a *Codat*. A codat may belong to only one peer group. If the same content must be published in two different peer groups, two different codats are created. The two codats may, of course, represent the same content.

Each codat is uniquely identified by a JXTA ID. All codats make their existence known to peer members by publishing a content advertisement.

A codat instance is a copy of a codat. Each codat copy may be replicated on different peers in the peer group. Each copy has the same codat id and an identical content.

Replicating codats within a peer group helps to ensure that each item of content is more readily available. For example, if an item has two instances residing on two different peers, only one of the peers needs to be alive to respond to the content request.

The JXTA protocols do not specify how codats are replicated. This decision is left to higher-level content service managers.

Part I

JXTA Core Specification

Chapter 1

IDs

1.1 Introduction

The JXTA protocols often need to refer to peers, peer groups, pipes and other JXTA resources. These references are presented in the protocols as JXTA IDs. JXTA IDs are a means for uniquely identifying specific peer groups, peers, pipes, codat and service instances. JXTA IDs provide unambiguous references to the various JXTA entities. There are six types of JXTA entities which have JXTA ID types defined: peergroups, peers, pipes, codats, module classes and module specifications. Additional JXTA ID types may be defined in the future.

JXTA IDs are normally presented as URNs. URNs are a form of URI that ‘... are intended to serve as persistent, location-independent, resource identifiers’. Like other forms of URI, JXTA IDs are presented as text. See *IETF RFC 2141* "[RFC2141](#)" for more information on URNs.

1.2 JXTA ID Properties

Every JXTA ID, regardless of format or type has the following properties:

- Unambiguous. It must be a complete reference to the resource.
- Relatively Unique. A JXTA ID must refer to a single resource.
- Canonical. References to the same resource should encode to the same JXTA ID. This enables JXTA IDs to be compared to determine if they refer to the same resource, but understandably may not be achievable by all ID Formats.
- Opaque. In their URN presentation JXTA IDs should be assumed to be opaque. The context of an ID within a protocol message generally is sufficient to establish its type. A JXTA binding may be able to interpret an ID if it supports the ID Format. Generally, only the immediate participants in a JXTA protocol need to understand the contents of a JXTA ID, if at all.

1.3 Using JXTA IDs in Protocols

When JXTA IDs are used within protocols they are manipulated as text string URIs. There are three operations available for URIs; compare, resolve, decompose. JXTA ID URIs are comparable for equality as strings. JXTA ID URIs can also be resolved to the resource they reference. Finally, JXTA ID URIs can optionally be decomposed and interpreted by JXTA bindings. In order to interpret a JXTA ID, a JXTA binding must support the *JXTA ID Format* used by that JXTA ID. For many JXTA protocols and operations it is not necessary to decompose the JXTA IDs.

1.4 Format of a JXTA ID URN

A JXTA ID is a standard URN in the JXTA ID namespace. JXTA ID URNs are identified by the URN namespace identifier `jxta`. Each JXTA ID URN also contains a JXTA ID Format keyword. The ID Format keyword indicates how the ID was created and may also allow JXTA bindings to extract additional information from the ID.

Two ID formats have been defined which are identified by the `jxta` and `uuid` keywords. It is possible to define additional JXTA ID Formats in order to refer to resources both within JXTA and to bridge to other technologies.

1.5 Example JXTA ID URNs

The following examples demonstrate valid JXTA ID presentation forms. These examples are not necessarily valid JXTA IDs.

Example 1.1 Sample JXTA IDs

- A. `urn:jxta:idform-1234567890`
 - B. `URN:jxta:idform-1234567890`
 - C. `urn:JXTA:idform-1234567890`
 - D. `urn:JXTA:IDForm-1234567890`
 - E. `urn:jxta:idform2-ABCDEFG`
 - F. `urn:jxta:idform3-31:08:66:42:67:::91:24::73`
-

In the preceding examples, A., B. and C. represent the same JXTA ID. Both the URN portion and the JXTA are case insensitive. Example D. is not equivalent to any of A., B. or C. because the data portion of the URN is case sensitive. In the six examples, four different JXTA ID Formats are used: 'idform', 'IDForm', 'idform2' and 'idform3'. Definition of ID Format names that differ only in character case is not recommended. The interpretation of the characters following the '-' is specific to each ID Format.

1.6 JXTA ID Representation

JXTA IDs are presented as URNs of the `jxta` namespace. The JXTA ID Namespace specifies additional restrictions upon the format of the URN. These requirements are detailed in [JXTA ID ABNF](#). The following figure uses the ABNF syntax as defined in *IETF RFC 2234* "[RFC2234](#)".


```

<JXTAURN>      ::= "urn:" <JXTANS> ":" <JXTAIDVAL>

<JXTANS>       ::= "jxta"

<JXTAIDVAL>    ::= <JXTAFMT> "-" <JXTAIDUNIQ>

<JXTAFMT>      ::= 1 * <URN chars>

<JXTAIDUNIQ>   ::= 1 * <URN chars>

<URN chars>    ::= <trans> | "%" <hex> <hex>

<trans>        ::= <upper> | <lower> | <number> | <other> |
                  <reserved>

<upper>        ::= "A" | "B" | "C" | "D" | "E" | "F" | "G" | "H" |
                  "I" | "J" | "K" | "L" | "M" | "N" | "O" | "P" |
                  "Q" | "R" | "S" | "T" | "U" | "V" | "W" | "X" |
                  "Y" | "Z"

<lower>        ::= "a" | "b" | "c" | "d" | "e" | "f" | "g" | "h" |
                  "i" | "j" | "k" | "l" | "m" | "n" | "o" | "p" |
                  "q" | "r" | "s" | "t" | "u" | "v" | "w" | "x" |
                  "y" | "z"

<hex>          ::= <number> | "A" | "B" | "C" | "D" | "E" | "F" |
                  "a" | "b" | "c" | "d" | "e" | "f"

<number>       ::= "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" |
                  "8" | "9"

<other>        ::= "(" | ")" | "+" | "," | "-" | "." |
                  ":" | "=" | "@" | ";" | "$" |
                  "_" | "!" | "*" | "/"

<reserved>     ::= "%" | "/" | "?" | "#"

```

Figure 1.1: JXTA ID ABNF

The jxta URN namespace does not currently define any special symbols from the `reserved` set.

1.7 JXTA ID Formats

JXTA IDs are designed to support multiple ID Formats. ID Formats allow JXTA developers to utilize existing naming and ID schemes within JXTA. In the JXTA ID presentation, the ID's 'Format' follows the JXTA URN namespace. Any JXTA ID Format which follows the general requirements for URNs and the [JXTA ID Properties](#) will be usable by conformant JXTA implementations.

1.8 JXTA ID Types

JXTA IDs may refer to many types of resources; pipes, peers, etc. Each JXTA ID format type may support references to one or more of these resource types. Currently, five standard resource types have been identified; peer groups, peers, pipes, content and service instances. Other types may be defined.

Each of the individual ID Types may provide additional requirements specific to its type.

1.8.1 Peer Group IDs

Peer Group IDs refer to peer groups. A peer group ID should canonically, uniquely and unambiguously refer to a peer group. Every ID Format must support this ID Type because all of the other ID Types refer to the peer group to which they belong. Every ID Format must support encoding of the World Peer Group. Support for other peer groups is optional. Example: You are defining an ID Format for Peer IDs based upon driver's license number. Driver's licenses are not organized into groups. This can be considered equivalent to all driver's licenses belonging to the same group, the global 'world peer group'.

1.8.2 Peer IDs

Peer IDs refer to peers. A Peer ID should canonically, uniquely and unambiguously refer to a peer. Support for this ID Type is optional. If a JXTA binding recognizes the ID Format, it should be able to determine the associated Peer Group ID from a Peer ID. This Peer Group ID identifies the peer group of which the peer is a member.

1.8.3 Codat IDs

Codat IDs refer to codats. A Codat ID should canonically, uniquely and unambiguously refer to a codat. Support for this ID Type is optional. If a JXTA binding recognizes the ID Format, it should be able to determine the associated Peer Group ID from a Codat ID. This Peer Group ID identifies the peer group of which the codat is a member.

1.8.4 Pipe IDs

Pipe IDs refer to pipes. A Pipe ID should canonically, uniquely and unambiguously refer to a pipe. Support for this ID Type is optional. If a JXTA binding recognizes the ID Format, it should be able to determine the associated Peer Group ID from a Pipe ID. This Peer Group ID identifies the peer group of which the pipe is a member.

1.8.5 Module Class IDs

A Module Class ID identifies a particular local behavior, that is, a specific API for each execution environment for which an implementation exists. A Module Class ID should canonically, uniquely and unambiguously refer to a module class as defined by an advertisement. Support for this ID Type is optional. If a JXTA binding recognizes the ID Type, it should be able to extract a Base Class ID from a Module Class ID. The Base Class ID allows applications to determine if two Module Class IDs differ only in the 'role' they perform. Module Spec ID's 'roles' allow for the same module to be reused within a group and have instances distinguished. This is necessary when, for example, a common database service is used, with each 'role' accessing a different data set.

1.8.6 Module Spec IDs

A ModuleSpecID uniquely identifies a particular network behavior (wire protocol and choreography) that may be embodied by a Jxta Module. There may be any number of implementations of a given Module Spec ID. A ModuleSpecID uniquely identifies an abstract module for which there may be multiple platform specific implementations. A ModuleSpecID is used to locate a compatible implementation such that it can be instantiated. All such implementations are assumed to be network compatible. A Module Spec ID should canonically, uniquely and unambiguously refer to a module specification. Support for this ID Type is optional. If a JXTA binding recognizes the ID Type, it should be able to extract a Module Class ID from a Module Spec ID.

1.9 JXTA ID Formats : 'jxta' ID Format

The 'jxta' ID Format is a required ID Format that is used for encoding 'well known' JXTA identifiers. All JXTA binding implementations must support this ID Format. There are three special reserved JXTA IDs; the Null ID, the World Peer Group ID and the Net Peer Group ID. The 'jxta' ID Format exists so that for these few 'well known' IDs only a single representation exists.

```
<JXTAJXTAURN>      ::= "urn:" <JXTANS> ":" <JXTAJXTAFMT> "-"  
                      <JXTAJXTAFMTID>  
  
<JXTAJXTAFMT>      ::= "jxta"  
  
<JXTAJXTAFMTID>    ::= <JXTANULL> | <JXTAWORLDDGROUP> | <JXTANETGROUP>  
  
<JXTANULL>         ::= "Null"  
  
<JXTAWORLDDGROUP>  ::= "WorldGroup"  
  
<JXTANETGROUP>     ::= "NetGroup"
```

Figure 1.2: JXTA ID : "jxta" ID Format ABNF

Chapter 2

Advertisements

2.1 Introduction

Advertisements are meta-data documents used by JXTA protocols to describe resources. Advertisements are used to describe peers, peer groups, pipes, content, services and many other types of resources. JXTA Advertisements are presented in XML. Many of the JXTA protocols depend on Advertisements to provide necessary information. JXTA protocols are used to pass Advertisements between peers.

Services can define new Advertisement types by sub-typing existing Advertisement types or by defining completely new Advertisements. Advertisement sub-types allow for additional information to be provided as well as richer meta-data.

Advertisements are composed of a series of hierarchically arranged elements. The elements may appear in any order within the advertisement. Each element can contain its data or additional elements. An element can also have attributes. Attributes are name-value string pairs. An attribute is used to store meta-data, which helps to describe the data within the element.

The Core JXTA Protocols rely on the following advertisements:

- [Peer Advertisement](#)
- [Peer Group Advertisement](#)
- [Module Class Advertisement](#)
- [Module Specification Advertisement](#)
- [Module Implementation Advertisement](#)

2.2 XML and JXTA Advertisements

All JXTA advertisements are represented in XML. XML provides a powerful means of representing data and metadata throughout a distributed system. XML provides a universal (software-platform neutral) representation:

- XML is programming language agnostic
 - XML is self-describing
 - XML content can be strongly-typed
 - XML ensures correct syntax
-

These advantages allow peers to manage and use Advertisements safely and to be able to ensure correct interactions with other peers. The Advertisements defined by the *JXTA Core Specification* and the *JXTA Standard Services* are specified using the *XML Schema Definition Language* "XSD2001-1" "XSD2001-2". Use of XML Schemas allows the advertisement contents to be strongly type-checked and semantically validated beyond the syntactical validation provided by XML with DTDs. Service and protocol authors are recommended to specify their Advertisements or Advertisement sub-types using XML Schema Language. DTDs are normally prepared from the schema descriptions for use in environments which do not support XML schema.

The other powerful feature of XML is its ability to be translated into other encodings such as HTML and WML. This feature allows peers that do not support XML to access advertised resources.

```
<xs:simpleType name="JXTAID">
  <xs:restriction base="xs:anyURI">
    <xs:pattern value="([uU][rR][nN]:[jJ][xX][tT][aA]:)\.\+\-\.\+"/>
  </xs:restriction>
</xs:simpleType>

<xs:complexType name="serviceParam">
  <xs:sequence>
    <xs:element name="MCID" type="jxta:JXTAID"/>
    <xs:element name="Parm" type="xs:anyType"/>
  </xs:sequence>
</xs:complexType>

<xs:complexType name="Cred">
  <xs:all>
  </xs:all>
</xs:complexType>
```

Figure 2.1: Common Advertisement Fragments Schemas

2.3 Peer Advertisement

A Peer Advertisement describes a peer and the resources it provides to the group. The Peer Advertisement holds specific information about the peer such as its unique id, its group id and optionally its name and descriptive information. It may also contain endpoint addresses and any run-time attributes that individual peer services want to publish (such as being a rendezvous peer for a group).

```
<xs:element name="PA" type="jxta:PA"/>
<xs:complexType name="PA">
  <xs:sequence>
    <xs:element name="PID" type="JXTAID"/>
    <xs:element name="GID" type="JXTAID"/>
    <xs:element name="Name" type="xs:string" minOccurs="0"/>
    <xs:element name="Desc" type="xs:anyType" minOccurs="0"/>
    <xs:element name="Svc" type="jxta:serviceParams" minOccurs="0" maxOccurs="unbounded"/>
  </xs:sequence>
</xs:complexType>
```

Figure 2.2: Peer Advertisement Schema

<PID>

This is a required element that uniquely identifies this peer. Each peer has a unique id.

<GID>

This is a required element that identifies the Peer Group to which this peer belongs.

<Name>

This is an optional string that can be associated with a peer. The name is not required to be unique unless the name is obtained from a centralized naming service that guarantees name uniqueness.

<Desc>

This is an optional string that can be used to index and search for a peer. The string is not guaranteed to be unique. Two peers may have the same keywords. The keywords string may contain spaces.

<Svc>

Any number of such elements may exist. Each of them describes the association between a group service denoted by its Class ID (the value of an MCID element), and arbitrary parameters encapsulated in a Parm element. For example, all accessible endpoint addresses are published in association with the Endpoint Service Class ID. The TLS Root certificate is published under the PeerGroup Class ID (There is a class ID for Peer Group as well). The flag that denotes that this peer is a rendezvous for this group is published under the Rendezvous Service Class ID. Ultimately, each service is responsible for what is published under its Class ID. The Service section may also optionally contain an element "isOff" meaning that this service is disabled. This element is used to convey a configuration choice made by the owner of the peer.

2.4 Peer Group Advertisement

A Peer Group Advertisement describes peergroup specific resources: name, group id, description, specification, and service parameters.

```
<xs:element name="PGA" type="jxta:PGA"/>

<xs:complexType name="PGA">
  <xs:sequence>
    <xs:element name="GID" type="jxta:JXTAID"/>
    <xs:element name="MSID" type="jxta:JXTAID"/>
    <xs:element name="Name" type="xs:string" minOccurs="0"/>
    <xs:element name="Desc" type="xs:anyType" minOccurs="0"/>
    <xs:element name="Svc" type="jxta:serviceParam" minOccurs="0" maxOccurs="unbounded"/>
  </xs:sequence>
</xs:complexType>
```

Figure 2.3: Peer Group Advertisement Schema

<GID>

This element provides the Peer Group ID. The Peer Group ID is the canonical way of referring to a group and uniquely identifies the peer group. See [Peer Group IDs](#) for more information on peer group ids.

<MSID>

Peer group Specification ID. This designates the module that provides the peer group mechanism itself for that group. The spec ID designates an abstraction of that mechanism. This abstraction is optionally described by a ModuleSpecAdvertisement, and any number of implementations may exist, each described by a ModuleImplAdvertisement. These advertisements may all be searched by this SpecID.

<Name>

This is an optional name that can be associated with a peergroup. The name is not required to be unique unless the name is obtained from a centralized naming service that guarantee name uniqueness.

<Desc>

This is an optional element provides descriptive information that may be used to index and search for a peergroup. The content of this element may not be unique. For example, two peergroups may have the same keywords.

<Svc>

Any number of such elements may exist. Each of them describes the association between a group service denoted by its Class ID (the value of an MCID element), and arbitrary parameters encapsulated in a Parm element. This optional parameter may only be meaningful to some services. It is used to configure a service specifically in relation with its use by this group. For example, a simple membership service may find an encrypted password list there.

2.5 Module Class Advertisement

A Module Class Advertisement describes a class of modules. That is, an expected local behavior and the expected API for each JXTA binding (that supports such modules). The purpose of this advertisement is to provide a description of what a particular Module Class ID stands for. A Module Class ID is what other modules or other code running on JXTA uses to designate modules which it depends upon. The ModuleClassAdvertisement is not required to provide a completely formal description of the module's behavior and API. It is intended for humans who want to create modules with a similar functionality. It is not required to publish a Module Class Advertisement for a Module Class ID to be valid, although it is a good practice.

```
<xs:element name="MCA" type="jxta:MCA"/>

<xs:complexType name="MCA">
  <xs:sequence>
    <xs:element name="MCID" type="jxta:JXTAID"/>
    <xs:element name="Name" type="xs:string" minOccurs="0"/>
    <xs:element name="Desc" type="xs:anyType" minOccurs="0"/>
  </xs:sequence>
</xs:complexType>
```

Figure 2.4: Module Class Advertisement Schema

<MCID>

Module Class ID. This is a required element that uniquely identifies the class. Each module class has a unique id. The class id representation is given in the Id Chapter.

<Name>

This is an optional name that can be associated with a class. The name is not required to be unique unless the name is obtained from a centralized naming service that guarantee name uniqueness.

<Desc>

Description. This is an optional string that can be used to describe and search for a class.

2.6 Module Specification Advertisement

A Module Specification Advertisement describes the specification of a module. That is, an expected on-wire behavior and protocol. The purpose of this advertisement is to provide a description of what a particular Module Specification ID stands for. A Module Specification ID is what other modules or other code running on JXTA uses to designate a particular network-compatible family of implementations of a given class. It is more importantly how a group implementation may designate the components which provide the various services that this group supports. All the built-in core peer group services (discovery, membership, resolver,...) are modules.

It is not required to publish a Module Spec Advertisement for a Module Spec ID to be valid, although it is a good practice.

A Module Spec Advertisement may also describe how to invoke and use a module. A Module may be used through its API, by locating an implementation, loading it and starting it, or a module may be usable through a pipe or through a proxy module. Modules which permit this include one or both of a Pipe Advertisement or the Module Spec ID of a proxy module, in their ModuleSpecID. Publication of the Module Spec Advertisement is of course required in that case.

A Module Specification Advertisement is not required to provide a completely formal description of the module's network behavior or protocol, it is intended for humans who want to create compatible implementation of that specification.

```
<xs:element name="MSA" type="jxta:MSA"/>

<xs:complexType name="MSA">
  <xs:sequence>
    <xs:element name="MSID" type="jxta:JXTAID"/>
    <xs:element name="Vers" type="xs:string"/>
    <xs:element name="Name" type="xs:string" minOccurs="0"/>
    <xs:element name="Desc" type="xs:anyType" minOccurs="0"/>
    <xs:element name="Crtr" type="xs:string" minOccurs="0"/>
    <xs:element name="SURI" type="xs:anyURI" minOccurs="0"/>
    <xs:element name="Parm" type="xs:anyType" minOccurs="0"/>
    <xs:element ref="jxta:PipeAdvertisement" minOccurs="0"/>
    <xs:element name="Proxy" type="xs:anyURI" minOccurs="0"/>
    <xs:element name="Auth" type="jxta:JXTAID" minOccurs="0"/>
  </xs:sequence>
</xs:complexType>
```

Figure 2.5: Module Specification Advertisement Schema

<MSID>

ModuleSpecID. This is a required element that uniquely identifies the specification. Each module specification has a unique id. The spec id representation is given in the Id Chapter.

<Vers>

The mandatory version of the specification that this advertises.

<Name>

This is an optional name that can be associated with a spec. The name is not required to be unique unless the name is obtained from a centralized naming service that guarantee name uniqueness.

<Desc>

Description. This is an optional string that can be used to describe and search for a spec.

<CRTR>

Creator. This optional element designates the creator of this specification.

<SURI>

Spec URI. This optional element is a URI that permits to retrieve a document containing the specification that this advertises.

<Parm>

Arbitrary parameters to be interpreted by each implementation.

<jxta:PipeAdvertisement>

Pipe advertisement. A pipe advertisement which this module binds to an input pipe and which thus may be used to establish a pipe to a nearby running implementation of this specification. Note that the element name is identical to the Pipe Advertisement document type since the entire element is an embedded pipe advertisement document.

<Proxy>

Proxy Spec ID. Optional ModuleSpecID of a proxy module that may be used in order to communicate with modules of this specification. Note that the process may be recursive. The proxy module may be usable via pipes, or through a subsequent proxy module, and itself require a subsequent authenticator. However publishers of modules should probably avoid such designs.

<Auth>

Authenticator Spec ID. Optional ModuleSpecID of an authenticator module that may be required in order to communicate with modules of this specification. Note that the process may be recursive. The authenticator module may be usable via pipes, or through a subsequent proxy module, and itself require a subsequent authenticator. However publishers of modules should probably avoid such designs.

2.7 Module Implementation Advertisement

A Module Implementation Advertisement describes one of the implementations of a module specification. Implementations of a given specification may be searched by the SpecID. An implementation may be selected by the type of environment in which it can be used (its compatibility statement) as well as by its name, description or the content of its parameters section.

A Module Implementation Advertisement also provides a means to retrieve all the necessary data required in order to execute the implementation being described. This information is encapsulated in the Code and PURI elements. The interpretation of these elements are subject to the the module's compatibility. For example, the standard peer group implementation of the Java reference implementation expects the <Code>element to specify a fully qualified Java class name that designates a subclass of `net.jxta.platform.Module` and PURI to be the URI of a downloadable package (a `.jarfile`). Other execution environments could expect the code to be inline within the <Code>element or even offer several options.

```
<xs:element name="MIA" type="jxta:MIA"/>

<xs:complexType name="MIA">
  <xs:sequence>
    <xs:element name="MSID" type="jxta:JXTAID"/>
    <xs:element name="Comp" type="xs:anyType"/>
    <xs:element name="Code" type="xs:anyType"/>
    <xs:element name="PURI" type="xs:anyURI" minOccurs="0"/>
    <xs:element name="Prov" type="xs:string" minOccurs="0"/>
    <xs:element name="Desc" type="xs:anyType" minOccurs="0"/>
    <xs:element name="Parm" type="xs:anyType" minOccurs="0"/>
  </xs:sequence>
</xs:complexType>
```

Figure 2.6: Module Implementation Advertisement Schema

<MSID>

ModuleSpecID. This is a required element that uniquely identifies the specification being implemented. The SpecID representation is given in the Id Chapter.

<Comp>

Compatibility. A mandatory arbitrary element that describes the environment in with this implementation may be executed. Each framework capable of loading and executing module has its own requirement on the contents of this element.

<Code>

This arbitrary element contains anything that is needed in addition to the package in order to load and execute the code of this implementation. In the case of a java implementation it contains the fully qualified class name containing the module's entry points. In other cases it may contain the entire code.

<PURI>

Package URI. This optional element is a URI that permits to retrieve a package containing the code of this implementation.

<Prov>

Provider. The provider of that implementation.

<Desc>

Description. This is an optional string that can be used to describe and search for a spec.

<Parm>

Parameter. Arbitrary parameters to be interpreted by the implementation's code.

Chapter 3

JXTA Core Protocols

3.1 Peer Resolver Protocol

3.1.1 Introduction

The *Peer Resolver Protocol* (PRP) provides a generic query/response interface applications and services can use for building resolution services. The PRP provides the ability to issue queries within a peer group and later identifying matching responses. Each query or response is addressed to a specific named handler. The Resolver Service of each peer group cooperates with the named handlers to provide the query routing strategies and policies. The named handlers provide the specific semantics for how the query is distributed and resolved within the peer group and how responses to the query are handled. In most situations the service or application has the best knowledge of the group topology and how the query should be routed. A query may be received and processed by any number of peers within the peer group, possibly all, and it is processed accordingly by a handler if such a named handler is registered on that peer.

Peers may also participate in the Shared Resource Distributed Index (SRDI). SRDI provides a generic mechanism JXTA services and applications can utilize a distributed index of shared resources with other peers that are grouped as a set of more capable peers such as rendezvous peers. These indices can be used to direct queries in the direction where the query is most likely to be answered, and repropagate messages to peers interested in these propagated messages.

3.1.2 Resolver Query Message

The resolver query message is used to send a resolver query to the named handler on one or more peers that are members of the peer group. The resolver query is sent as a query string to a specific handler. Each query has a unique Id. The query string can be any string that will be interpreted by the targeted handler.

```
<xs:element name="ResolverQuery" type="jxta:ResolverQuery"/>

<xs:complexType name="ResolverQuery">
  <xs:sequence>
    <xs:element ref="jxta:Cred" minOccurs="0" />
    <xs:element name="SrcPeerID" type="jxta:JXTAID" />
    <!-- This could be extended with a pattern restriction -->
    <xs:element name="HandlerName" type="xs:string" />
    <xs:element name="QueryID" type="xs:string" />
    <xs:element name="HC" type="xs:unsignedInt" />
    <!-- For historical reasons the query is a whole flattened document -->
    <xs:element name="Query" type="xs:anyType" />
  </xs:sequence>
</xs:complexType>
```

Figure 3.1: Resolver Query Schema

<jxta:Cred>

The credential of the sender.

<HandlerName>

A string that specifies the destination of this query.

<SrcPeerID>

The id of the peer originating the query (as a URN).

<QueryID>

An opaque identifier to be used by the querier to match replies. The **<QueryID>** should be included in the responses to this query.

<HC>

specifies the number of hops the query has been through The **<HC>** should be incremented by each peer that forwards the query.

<Query>

Contains the query.

Example 3.1 Resolver Query

```

<?xml version="1.0"?>

<!DOCTYPE jxta:ResolverQuery>

<jxta:ResolverQuery xmlns:jxta="http://jxta.org">
  <HandlerName>
    urn:jxta:uuid-DEADBEEFDEAFBABAEEEDBABE0000000305
  </HandlerName>
  <jxta:Cred>
    JXTACRED
  </jxta:Cred>
  <QueryID>
    0
  </QueryID>
  <HC>
    0
  </HC>
  <SrcPeerID>
    urn:jxta:uuid-59616261646162614A7874615032503304BD268FA4764960AB93A53D7F15044503
  </SrcPeerID>
  <Query>
    &lt;?xml version="1.0"?>

    &lt;!DOCTYPE jxta:DiscoveryQuery>

    &lt;jxta:DiscoveryQuery xmlns:jxta="http://jxta.org">
      &lt;Type>
        0
      &lt;/Type>
      &lt;Threshold>
        50
      &lt;/Threshold>
      &lt;PeerAdv>
        &lt;?xml version="1.0"?>

        &lt;!DOCTYPE jxta:PA>
          ... remainder omitted for brevity ...
        &lt;/jxta:PA>
      &lt;/PeerAdv>
      &lt;Attr>
        &lt;/Attr>
        &lt;Value>
          &lt;/Value>
        &lt;/jxta:DiscoveryQuery>
      </Query>
    </jxta:ResolverQuery>

```

3.1.3 Resolver Response Message

A resolver response message is used to send a response to a resolver query message.

```
<xs:element name="ResolverResponse" type="ResolverResponse"/>

<xs:complexType name="ResolverResponse">
  <xs:sequence>
    <xs:element ref="jxta:Cred" minOccurs="0"/>
    <xs:element name="ResPeerID" type="jxta:JXTAID" minOccurs="0" />
    <xs:element name="HandlerName" type="xs:string"/>
    <xs:element name="QueryID" type="xs:string"/>
    <!-- For historical reasons the response is a whole flattened document -->
    <xs:element name="Response" type="xs:anyType"/>
  </xs:sequence>
</xs:complexType>
```

Figure 3.2: Resolver Response Schema

- <jxta:Cred>**
The credential of the respondent.
- <HandlerName>**
Specifies how to handle the response.
- <ResPeerID>**
The id of the peer originating the response (as a URN).
- <QueryID>**
The query identifier of the query to which this is a response.
- <Response>**
The responses.

Example 3.2 Resolver Response

```

<?xml version="1.0"?>

<!DOCTYPE jxta:ResolverResponse>

<jxta:ResolverResponse xmlns:jxta="http://jxta.org">
  <HandlerName>
    urn:jxta:uuid-DEADBEEFDEAFBABAFEEDBABE0000000305
  </HandlerName>
  <jxta:Cred>
    JXTACRED
  </jxta:Cred>
  <QueryID>
    0
  </QueryID>
  <ResPeerID>
    urn:jxta:uuid-59616261646162614A7874615032503304BD268FA4764960AB93A53D7F15044503
  </ResPeerID>
  <Response>
    &lt;?xml version="1.0"?>

    &lt;!--DOCTYPE jxta:DiscoveryResponse-->

    &lt;jxta:DiscoveryResponse xmlns:jxta="http://jxta.org">
      &lt;Count>
        1
      &lt;/Count>
      &lt;Type>
        2
      &lt;/Type>
      &lt;PeerAdv>
        &amp;lt;?xml version="1.0"?>

        &amp;lt;!--DOCTYPE jxta:PA-->

        &amp;lt;jxta:PA xmlns:jxta="http://jxta.org">
          ... remainder omitted for brevity ...
        &amp;lt;/jxta:PA>
      &lt;/PeerAdv>
      &lt;Response Expiration="7200000">
        &amp;lt;?xml version="1.0"?>

        &amp;lt;!--DOCTYPE jxta:PipeAdvertisement-->

        &amp;lt;jxta:PipeAdvertisement xmlns:jxta="http://jxta.org">
          &amp;lt;Id>
            urn:jxta:uuid-59616261646162614 ↵
            E50472050325033D1D1D1D1D1D1D1D1D1D1D1D1D1D104
          &amp;lt;/Id>
          &amp;lt;Type>
            JxtaPropagate
          &amp;lt;/Type>
          &amp;lt;Name>
            JxtaTalkUserName.IP2PGRP
          &amp;lt;/Name>
        &amp;lt;/jxta:PipeAdvertisement>
      &lt;/Response>
    &lt;/jxta:DiscoveryResponse>
  </Response>
</jxta:ResolverResponse>

```

3.1.4 Resolver SRDI Message

The resolver SRDI message is used to send a resolver SRDI message to the named handler on one or more peers that are members of the peer group. The resolver SRDI message is sent to a specific handler. The payload string may be any string that will be interpreted by the targeted handler.

```
<xs:element name="ResolverSRDI" type="jxta:ResolverSRDI"/>

<xs:complexType name="ResolverSRDI">
  <xs:all>
    <xs:element name="HandlerName" type="xs:string"/>
    <xs:element ref="jxta:Cred" minOccurs="0"/>
    <xs:element name="Payload" type="xs:anyType"/>
  </xs:all>
</xs:complexType>
```

Figure 3.3: Resolver SRDI Schema

<HandlerName>

A string that specifies the destination of this message.

<jxta:Cred>

The credential of the sender.

<Payload>

Contains the payload.

Example 3.3 Resolver SRDI

```
<?xml version="1.0"?>

<!DOCTYPE jxta:ResolverSRDI>

<jxta:ResolverSRDI xmlns:jxta="http://jxta.org">
  <HandlerName>
    urn:jxta:uuid-DEADBEEFDEAFBABAFEEDBABE0000000305
  </HandlerName>
  <jxta:Cred>
    JXTACRED
  </jxta:Cred>
  <Payload>
    &lt;?xml version="1.0"?>

    &lt;!DOCTYPE jxta:GenSRDI>
      ... remainder omitted for brevity ...
    &lt;/jxta:GenSRDI>
  </Payload>
</jxta:ResolverSRDI>
```

3.1.5 Listener and Element Naming

The PRP communicates by exchanging Messages using the Endpoint Service. Endpoint Addresses specify a handler name. The PRP attaches a listener by that name to the [Endpoint Service](#).

All PRP implementations must use the same scheme for building their handler names. The convention used by all services of the world peer group is to use the concatenation of the service name, the peer group ID, and a value unique within the service.


```

<JXTARSLVRRSQRY> ::= <JXTARSLVRNAM> <JXTAIDVAL> <JXTARSLVRQRYTAG>
<JXTARSLVRRSRSP> ::= <JXTARSLVRNAM> <JXTAIDVAL> <JXTARSLVRRSPTAG>
<JXTARSLVRRSSRDI> ::= <JXTARSLVRNAM> <JXTAIDVAL> <JXTARSLVRSRDITAG>
<JXTARSLVRQRYTAG> ::= "ORes"
<JXTARSLVRRSPTAG> ::= "IRes"
<JXTARSLVRSRDITAG> ::= "Isrdi"
<JXTARSLVRNAM> ::= "jxta.service.resolver"
<JXTAIDVAL> ::= See

```

Figure 3.4: Listener Naming Syntax ABNF

The listeners used by the PRP are named as follows:

Queries

jxta.service.resolver[group unique Id string]ORes (*ORes* is a literal string)

Responses

jxta.service.resolver[group unique Id string]IRes (*IRes* is a literal string)

Srди

jxta.service.resolver[group unique Id string]Isrди (*Isrди* is a literal string)

Query and response messages are included in messages as elements named as follows:

- *Queries*: [group unique Id string]ORes (*ORes* is a literal string)
- *Responses*: [group unique Id string]IRes (*IRes* is a literal string)
- *Srди*: [group unique Id string]Isrди (*Isrди* is a literal string)

3.1.6 Behaviour

3.1.6.1 Handler Name

The Handler Name in PRP messages plays a role similar to that of the handler name in the Endpoint Message addresses: it is a demultiplexing key that specifies how, by which higher-level protocol, or by which module, the message is to be processed.

In the Java and 'C' reference implementations, the users of the PRP are typically services. Each instance of a given service (one per peer per group that uses this service) generates a handler name that is unique on its peer, but will be identical for the instances of this service on other peers. This is by convention achieved by concatenating the service name (which is unique in the group), the group id, which is unique in the peer, and a additional parameter which serves to discriminate between several handlers used by the same service, if needed.

The handler name is used both to register the appropriate handler for incoming queries or responses, and as a destination for outgoing queries or responses. For convenience, most clients of the resolver do define two names: one for propagated messages (mostly queries), and one for unicast messages (mostly responses).

The PRP typically uses the Rendezvous Service to disseminate a query to multiple peers or unicast messages to send queries to specified peers.

The PRP should refuse, and the existing reference implementations shall refuse the registration of more than one handler with the same name. A service should register for any handler name that it uses as a destination, thereby preventing other services from

registering themselves to receive these messages. This means that in principle a service or application that receives queries or responses from a service instance on another peer is de-facto the local instance of that service and should handle these messages as specified. PRP is designed for same-to-same communication, not client-server.

3.1.6.2 Policies and Quality of Service

The PRP does not guarantee peers that define a query handler name will receive that query, nor does it mandate that all peers that define this handler name will receive it. Only a best effort is made at disseminating the query in a way that maximizes the chance of obtaining a response, if one can be obtained.

There is no guarantee that a response to a resolver query request will be made. It is important to point that response to a ResolverQuery request is optional. A peer is not required to respond.

There is no guarantee that a Resolver SRDI Message will be honored. It is important to point out that accepting a Resolver SRDI Message is optional. A peer is not required to accept the message.

The PRP does not assume the presence of reliable message delivery. Multiple Resolver query messages may be sent--none, one, multiple or redundant responses may be received.

The PRP provides a generic mechanism for services to send queries, and receive responses and SRDI messages. As a service, the reference implementation helps other services by taking care of all messaging aspects, caching queries responses, and SRDI messages and in forwarding queries, based on the invoker's decision. The PRP performs authentication, and verification of credentials and drops incorrect messages.

The actual task of propagating a query to the next set of peers is delegated to the [Rendezvous Protocol](#). The Rendezvous service is responsible for determining the set of peers that should receive a message being propagated, but never automatically re-propagates an incoming propagated message. It is left to the service (query handler) handling the message to determine if further propagation should be performed. The PRP's policy is the following: if the query handler does not instruct the PRP to discard the query, and if the local peer is a rendezvous, then the query is re-propagated (within the limits of loop and TTL rules enforced by the Rendezvous service). In addition, if instructed by the query handler, an identical query may be issued with the local peer as the originator.

3.2 Endpoint Routing Protocol

The JXTA network is ad hoc, multi-hop, and adaptive by nature. Connections in the network may be transient, and message routing is nondeterministic. Routes may be unidirectional and change rapidly. Peers may join and leave frequently. A peer inside a firewall can send a message directly to a peer outside a firewall. But a peer outside the firewall cannot establish a connection directly with a peer inside the firewall.

The Endpoint Routing Protocol (ERP) defines a set of request/query messages that are processed by a routing service to help a peer route messages to their destination.

When a peer is asked to send a message to a given peer endpoint address, it looks in its local cache to see if it has a route to this peer. If it does not find a route, it sends a route resolver query message to its available peer routers asking for routing information. A peer can have as many peer routers as it can find or they can be pre-configured. Pre-configured routers are optional.

The peer routers provide the low-level infrastructure to route messages between two peers in the network. Any number of peers in a peer group can elect themselves to become peer routers for other peers. Peer routers offer the ability to cache route information, as well as bridging different physical or logical networks. A peer can dynamically find its router peer via a qualified discovery search. A peer can find out if a peer it has discovered is a peer router via the peer advertisement <Parms> element.

When a peer router receives a route query, if it knows a route to the destination, it answers the query by returning the route information as an enumeration of hops. Once a route has been discovered, a message can be sent to the first router and that router will use the route information to route the message to the destination peer. The route is ordered from the next hop to the final destination peer. At any point the routing information may become obsolete requiring the current router to discover a new route in order to complete the message delivery.

The peer endpoint adds extra routing information to the messages sent by a peer. When a message goes through a peer, the endpoint of that peer leaves its trace on the message. The trace can be used for loop detection, and to discard recurrent messages. The trace is also used to record new route information by peer routers.

- The <RECIPIENT> and <RECIPIENTPARAM> separated by a /.
- The <RECIPIENT>.

3.2.2 Route information

Route information is represented as follow:

```
<xs:element name="APA" type="jxta:APA"/>

<xs:complexType name="jxta:APA">
  <xs:sequence>
    <xs:element name="PID" minOccurs="0" type="jxta:JXTAID"/>
    <xs:element name="EA" type="xs:anyURI" minOccurs="0" maxOccurs="unbounded"/>
  </xs:sequence>
</xs:complexType>

<xs:element name="RA" type="jxta:RA"/>

<xs:complexType name="jxta:RA">
  <xs:sequence>
    <xs:element name="DstPID" minOccurs="0" type="jxta:JXTAID"/>
    <xs:element name="Dst">
      <xs:sequence>
        <xs:element ref="jxta:APA" maxOccurs="1"/>
      </xs:sequence>
    </xs:element>
    <xs:element name="Hops" minOccurs="0">
      <xs:sequence>
        <xs:element ref="jxta:APA" maxOccurs="unbounded"/>
      </xs:sequence>
    </xs:element>
  </xs:sequence>
</xs:complexType>
```

Figure 3.6: Route Advertisement

<PID>

The Peer ID of the peer described by this advertisement. When not included the advertisement is used in a context where the Peer ID is already known, such as in the <Dst> element of a Route Advertisement.

<EA>

An Endpoint Address for the peer.

<DstPID>

The Peer ID of the peer described by this advertisement. When not included the advertisement is used in a context where the Peer ID is already known, such as in a Peer Advertisement.

<Dst>

Contains an Access Point Advertisement containing a list of endpoint addresses associated with the specified peer id.

<Hops>

A semi-ordered collection of Access Point Advertisements describing a route to the peer indicated by <DstPID>. The individual steps of the route are co-mingled with routing alternatives. As the route is traversed peers route messages by forwarding to the earliest peer in the list that they can reach.

The time-to-live parameter is measured in hops and specifies how long this route is valid. The creator of the route can decide how long this route will be valid. The gateways are defined as an ordered sequence of peer IDs which define the route from the source peer to the destination peer. The sequence may not be complete, but at least the first gateway should be present. The first gateway is sufficient to initially route the messages. The remaining gateway sequence is optional.

Peer routers will typically cache route information. Any peer can query a peer router for route information. Any peer in a peer group may become a peer router.

3.2.3 Route Query Message

This message is sent by peers to request route information for another peer. Route Query Messages are transmitted as queries within [Resolver Query Message](#).

```
<xs:element name="ERQ" type="jxta:ERQ"/>

<xs:complexType name="jxta:ERQ">
  <xs:sequence>
    <xs:element name="Dst" type="jxta:JXTAID"/>
    <xs:element name="Src">
      <xs:element ref="jxta:RA"/>
    </xs:element>
  </xs:sequence>
</xs:complexType>
```

Figure 3.7: Endpoint Router Query

<Dst>

The peer id of the peer who's route is desired.

<Src>

Route advertisement of the peer requesting route information. This route information is needed to ensure there is a return route for responses.

3.2.4 Route Response Message

This message is sent by peers in response to Route Query Messages. The Route Response Message contains a route advertisement for the destination peer. Route Response Messages are transmitted as responses within [Resolver Response Message](#).

```
<xs:element name="ERR" type="jxta:ERR"/>

<xs:complexType name="jxta:ERR">
  <xs:sequence>
    <xs:element name="Dst">
      <xs:element ref="jxta:RA"/>
    </xs:element>
    <xs:element name="Src">
      <xs:element ref="jxta:RA"/>
    </xs:element>
  </xs:sequence>
</xs:complexType>
```

Figure 3.8: Endpoint Router Response Message

<Dst>

The Route Advertisement for the peer which was requested in the Endpoint Router Query.

<Src>

Route Advertiement for the destination peer.

Chapter 4

Core JXTA Message Transport Bindings

4.1 Endpoint Service

4.1.1 Description

The Endpoint Service is responsible for performing end-to-end messaging between two JXTA peers, using one of the underlying JXTA transport protocols, such as the JXTA TCP or HTTP bindings.

The Endpoint Service is primarily used by other services or applications that need to have an understanding of the network topology, such as the Resolver Service or the Propagation Service.

The Endpoint Service is not responsible for routing messages for peers that are not directly connected to each other. This task is performed by the [Endpoint Router Transport Protocol](#) which provides the illusion that the source and destination peers are directly connected.

4.1.2 Protocol

When the Endpoint Service transmits a message it may add a single element to the message: the source peer ID.

The element name is: `jxta:EndpointHeaderSrcPeer` and its value is a textual UTF-8 representation of the peer ID at the point of emission of the message. This information is optional and is used by the emitter endpoint service itself to detect and eliminate propagated messages that loop back to the emitter.

If this element is not present the message is assumed to not be looping back.

The endpoint service expects incoming and outgoing messages to have a source address and a destination address. The encapsulation of that information is specified by the message wire format being used.

Two additional common elements are optionally used by Message Transports in conjunction with the Endpoint Service. Message Transports may typically provide received messages to the Endpoint Service containing these elements and the Endpoint service will dispatch the messages based upon their content. Message Transports may use alternate mechanisms for determining message source and destination addresses and need not use these elements.

The `jxta:EndpointSourceAddress` contains an Endpoint Address for the source of this message. The source address has a variety of meanings based upon the usage of the underlying Message Transport. For low level transports such as TCP or HTTP the source address is the return address of the peer from which the message was received, ie. the hop address. For higher level Message Transports such as the Endpoint Router Transport or the TLS transport the source address is the virtual Endpoint Address of the peer which originated the message regardless of any intervening hops the message may have made.

The `jxta:EndpointDestinationAddress` contains an Endpoint Address which will be used by the Endpoint Service to dispatch a received message to the recipient specified by the service name and service parameter. The protocol address is also provided to the recipient service and can be used in some protocols for determining how the message was received. For example a service may wish to handle messages which were sent directly differently than messages which were sent via propagation.

4.2 Endpoint Router Transport Protocol

4.2.1 Description

The Endpoint Router is a logical JXTA Transport Protocol that sits below the Endpoint Service and beside the other Transport Protocols such as the JXTA TCP and HTTP Transport Protocols.

The Endpoint Router is responsible for exchanging messages between peers that do not have a direct connection between each other. The Endpoint Router provides a virtual direct connection to the peer's Endpoint Service.

4.2.2 Protocol

The Endpoint Router protocol defines a set of queries and responses used to communicate with instances of the Endpoint Router on other peers.

- **Route Query:** when the Endpoint Router is requested to send a message to a peer for which it does not have yet a route for, it sends a Route Query request to other peers. Peers that have an route for the given peer answers with Route Response.
- **Route Response:** a peer that desires inform another peer about a give route sends a Route Response to the peer. A Route Response is replied following up a Route Query.

In addition, the Endpoint Router defines an informational message that requires no reply.

- **NACK:** a NACK is sent by any peer that detects that a route used by another peer is not valid. Typically, this happens by a router peer that are requested to route a message to peer for which it does not have a route itself. NACK messages are optional: routers are not required to send them, and while a NACK is typically sent to the source peer of the message, peers can send NACK to other peers of their choice.

These messages are sent and received by the EndpointRouter using the JXTA Resolver Service.

4.2.3 EndpointRouter Message Element

The Endpoint Router Transport Protocol appends its own message element to each message it transports. The name of the message element is `JxtaEndpointRouter` and contains an XML document containing the following:


```

<xs:element name="ERM" type="jxta:ERM"/>

<xs:complexType name="ERM">
  <xs:sequence>
    <xs:element name="Src" type="jxta:JXTAID" />
    <xs:element name="Dest" type="xs:anyURI" />
    <xs:element name="LastHop" minOccurs="0" type="jxta:JXTAID" />
    <xs:element name="Fwd">
      <xs:complexType>
        <xs:sequence>
          <xs:element ref="jxta:APA" maxOccurs="unbounded" />
        </xs:sequence>
      </xs:complexType>
    </xs:element>
    <xs:element name="Rvs" minOccurs="0">
      <xs:complexType>
        <xs:sequence>
          <xs:element ref="jxta:APA" maxOccurs="unbounded" />
        </xs:sequence>
      </xs:complexType>
    </xs:element>
  </xs:sequence>
</xs:complexType>

```

Figure 4.1: Endpoint Router Message Element

<Src>

The peer id of the originator of this message.

<Dest>

The Endpoint Address of the intended final destination of this message. Includes final routing information of the service name and service parameter.

<LastHop>

The peer id of the most recent peer to forward this message. When a peer forwards a message, it should update this field with its own Endpoint Address.

<Fwd>

A loosely ordered list of Access Point Advertisements which describe a probable route for this message. Loosely ordered because it may describe alternative routes. In order to route the message peers should scan the list for the *last* peer for which they recognize and route message via that peer.

<Rvs>

An ordered array of Access Point Advertisements which describe the known reverse route for this message. This is the path which the message has traveled thus far. Due to unidirectional links this route may not be suitable for sending response messages.

4.2.3.1 EndpointRouter Endpoint Address format

Since the EndpointRouter is a transport protocol, it has its own Endpoint Address format, which is:

```
jxta://[PeerID unique value]
```

Figure 4.2: JXTA Endpoint Router Address Format

Chapter 5

Messages

5.1 Introduction

Messages are the basic unit of data exchange between peers. Pipes send and receive messages to and from services; any protocol implemented by a service will send and receive messages. Messages are encoded using ‘wire’ representations for transmission. Each JXTA transport will use the message representations most appropriate for its characteristics and the peers’ preferences. See [JXTA Message Wire Representations](#) for information about representations.

5.2 Message

A message is a set of named and typed contents called elements. Thus a message is essentially a set of name/value pairs. The content can be an arbitrary type. Many core services send XML advertisements as message element contents.

As a message passes down the JXTA protocol stack (applications, services, endpoint and transports), each level may add one or more named elements to the message. As a message is passed back up the stack on the receiving peer, the protocol handlers should remove those elements.

A message is an ordered sequence of message elements. The most recently added element appears at the end of the message.

5.3 Element

A message element contains a namespace, an optional name, an optional type, an optional signature or digest and content.

5.3.1 Namespace

Every element is assigned to a namespace. Namespaces are used to organize elements used by different message users and transports within the same message.

Two namespaces names are considered equivalent if their representation in canonical UTF8 (NFC) (see *Unicode Standard Annex #15 : Unicode Normalization Forms "USA15"*) is byte-for-byte identical.

Two message element namespaces are pre-defined, "" (empty string) and `jxta`. The "" namespace is reserved for user applications and services--none of the JXTA protocols or services will use or modify elements in this namespace.

The `jxta` namespace is reserved for internal use by the JXTA protocols and services. Applications should not create, manipulate or assume the interpretation of any of the content of elements in the `jxta` namespace. In some bindings, applications may be forbidden from accessing or creating elements in the `jxta` namespace.

Use of namespaces by services and applications other than the "" namespace is optional. Namespaces require no formal registration as the protocols used need only be agreed upon by the participants.

5.3.2 Name

Elements may have an optional name. Elements in the same message may have the same name.

5.3.3 Type

A type is specified as a MIME type. See ["RFC2046"](#). The MIME type is encoded in canonical UTF8 (NFC) using the presentation and encoding of *Multipurpose Internet Mail Extensions (MIME) Part Two: Media Types* ["RFC2046"](#), ie. any encoding specified by RFC 2046 is performed before the string is encoded into UTF8 from its native representation.

The type is used by the applications and services that process the element. There is no restriction on the set of MIME types that can be used by applications and services. In addition to the applications and services which use the particular element, the type of the element may also be examined by the JXTA message transport to determine how to format the message element to ensure the most efficient transfer.

If the type is not specified for an element "application/octet-stream" is assumed.

5.3.4 Content

The contents of the Element data are opaque to except to the applications and services which use these elements.

Part II

JXTA Standard Services

Chapter 6

Standard Protocols

6.1 Peer Discovery Protocol

6.1.1 Introduction

The Peer Discovery Protocol is used to discover any published peer resource. Resources are represented as advertisements. A resource can be a peer, a peergroup, a pipe, a module, or any resource that has an advertisement. Each resource must be represented by an advertisement.

The Peer Discovery Protocol (PDP) enables a peer to find advertisements in its group. The PDP protocol is the discovery protocol of the world peergroup. Custom discovery services may choose to leverage PDP. If a peer group does not need to define its own discovery protocol, it may use the world peergroup PDP.

The intent is for PDP to provide the essential discovery infrastructure for building and bootstrapping high-level discovery services. In many situation, discovery information is better known by a high-level service, because the service may have a better knowledge of the group topology.

The PDP protocol provides a basic mechanism to discover advertisements while providing hooks so high-level services and applications can participate in the discovery process. Services should be able to give hints to improve discovery (i.e. decide which advertisements are the most valuable to cache).

The PDP protocol utilizes the resolver protocol to route queries and responses.

6.1.2 Discovery Query Message

The discovery query message is used by peers to send discovery requests when searching for advertisements.

```

<xs:element name="DiscoveryQuery" type="jxta:DiscoveryQuery"/>

<xsd:simpleType name="DiscoveryQueryType">
  <xsd:restriction base="xsd:string">
    <!-- peer -->
    <xsd:enumeration value="0"/>
    <!-- group -->
    <xsd:enumeration value="1"/>
    <!-- adv -->
    <xsd:enumeration value="2"/>
  </xsd:restriction>
</xsd:simpleType>

<xs:complexType name="DiscoveryQuery">
  <xs:sequence>
    <xs:element name="Type" type="jxta:DiscoveryQueryType"/>
    <xs:element name="Threshold" type="xs:unsignedInt" minOccurs="0"/>
    <xs:element name="Attr" type="xs:string" minOccurs="0"/>
    <xs:element name="Value" type="xs:string" minOccurs="0"/>
    <!-- The following should refer to a peer adv, but is instead a whole doc for ↵
        historical reasons -->
    <xs:element name="PeerAdv" type="xs:string" minOccurs="0"/>
  </xs:sequence>
</xs:complexType>

```

Figure 6.1: Discovery Query Schema

<Type>

Only advertisements of requested type will be matched. Possible values are:

- 0
Peer Advertisements
- 1
Peergroup Advertisements
- 2
Any Advertisements

<Threshold>

specifies the maximum number of advertisements that each responding peer should provide. The total number of results received depends on the number of peers that respond and the advertisements they have. If <Type> is 0 (Peer Advertisements) and <Threshold> is 0, then the query has a special meaning: its objective is to collect Peer Advertisements of respondents. Therefore any peer should respond to such a query, even though no results are to be included.

<PeerAdv>

If present, the advertisement of the requestor.

<Attribute>, <Value>

Must either be both present or absent. If absent, then each respondent should supply a random set of advertisements of the appropriate type up to <Threshold> count. Only advertisements containing an element whose name matches <Attribute> and that also contains a value matching <Value> are eligible to be found. <Value> may begin or end with "*", or both. In that case <Value> will match all values that end with or beginning with, or contain the rest of the string. If <Value> contains only "*" the result is unspecified. Some implementations may choose not match any advertisement for <Value> "*".

Example 6.1 Discovery Query

```
<?xml version="1.0" encoding="UTF-8"?>

<jxta:DiscoveryQuery>
  <Type>2</Type>
  <Threshold>1</Threshold>
  <Attr>Name</Attr>
  <Value>*sidus*</Value>
  <PeerAdv>
    &lt;?xml version="1.0"?>

    &lt;!DOCTYPE jxta:PA>

    &lt;jxta:PA xmlns:jxta="http://jxta.org">
      &lt;PID>
        urn:jxta:uuid-59616261646162614 ↵
        A7874615032503304BD268FA4764960AB93A53D7F15044503
      &lt;/PID>
      ... remainder omitted for brevity ...
    &lt;/jxta:PA>
  </PeerAdv>
</jxta:DiscoveryQuery>
```

6.1.3 Discovery Response Message

A Discovery response message is used by a peer to respond to a discovery query message.

```

<xs:element name="DiscoveryResponse" type="jxta:DiscoveryResponse"/>

<xs:complexType name="DiscoveryResponse">
  <xs:sequence>
    <xs:element name="Type" type="jxta:DiscoveryQueryType"/>
    <xs:element name="Count" type="xs:unsignedInt" minOccurs="0"/>
    <xs:element name="Attr" type="xs:string" minOccurs="0"/>
    <xs:element name="Value" type="xs:string" minOccurs="0"/>
    <!-- The following should refer to a peer adv, but is instead a whole doc for ↵
         historical reasons -->
    <xs:element name="PeerAdv" minOccurs="0">
      <xs:complexType>
        <xs:simpleContent>
          <xs:extension base="xs:string">
            <xs:attribute name="Expiration" type="xs:unsignedLong"/>
          </xs:extension>
        </xs:simpleContent>
      </xs:complexType>
    </xs:entry>
    <xs:element name="Response" maxOccurs="unbounded">
      <xs:complexType>
        <xs:simpleContent>
          <xs:extension base="xs:string">
            <xs:attribute name="Expiration" type="xs:unsignedLong"/>
          </xs:extension>
        </xs:simpleContent>
      </xs:complexType>
    </xs:entry>
  </xs:sequence>
</xs:complexType>

```

Figure 6.2: Discovery Response Schema

<Type>

The type of all the advertisements returned in the <Response> element(s).

<Count>

If present, the number of <Response> element(s) included in this response.

<PeerAdv>

If present, the advertisement of the respondent. The Expiration attribute is the associated relative expiration time in milliseconds.

<Attribute>, <Value>

If present, reflects that of the `DiscoveryQuery` to which this is the response.

<Response>

An advertisement. The Expiration attribute is the associated relative expiration time in milliseconds.

Example 6.2 Discovery Response

```

<?xml version="1.0" encoding="UTF-8"?>

<jxta:DiscoveryResponse>
  <Type>2</Type>
  <Count>1</Count>
  <Attr>Name</Attr>
  <Value>*sidus*</Value>
  <PeerAdv>
    &lt;?xml version="1.0"?>

    &lt;!DOCTYPE jxta:PA>

    &lt;jxta:PA xmlns:jxta="http://jxta.org">
      &lt;PID>
        urn:jxta:uuid-59616261646162614 ↵
        A7874615032503304BD268FA4764960AB93A53D7F15044503
      &lt;/PID>
      ... omitted ...
    &lt;/jxta:PA>
  </PeerAdv>
  <Response Expiration="36000000">
    &lt;?xml version="1.0"?>

    &lt;!DOCTYPE jxta:PipeAdvertisement>

    &lt;jxta:PipeAdvertisement xmlns:jxta="http://jxta.org">
      &lt;Id>
        urn:jxta:uuid-094 ↵
        AB61B99C14AB694D5BFD56C66E512FF7980EA1E6F4C238A26BB362B34D1F104
      &lt;/Id>
      &lt;Type>
        JxtaUnicastSecure
      &lt;/Type>
      &lt;Name>
        JxtaTalkUserName.sidus
      &lt;/Name>
    &lt;/jxta:PipeAdvertisement>
  </Response>
</jxta:DiscoveryResponse>

```

6.1.4 Behaviour

The PDP does not guarantee peers that receive a query will respond to the query, nor does it mandate that the number of advertisements requested will be honored. Only a best effort is made at matching the query to results in the respondent's cache.

There is no guarantee that a response to a discovery query request will be made. It is important to point out that responding to a DiscoveryQuery request is optional. A peer is not required to respond to a DiscoveryQuery request.

A reliable transport is optional with the PDP. Multiple Discovery query messages may be sent. None, one, multiple or redundant responses may be received.

A peer may receive a DiscoveryResponse that is not a response to any DiscoveryQuery initiated by the peer, this mechanism provides the ability to remote publish a resource.

The PDP provides a mechanism for services to query the network for JXTA resources, and receive responses. As a service, the reference implementation helps other services by taking care of all messaging aspects, caching, and expiring advertisements.

The actual task of propagating, and re-propagating a query to the next set of peers is delegated to the Resolver Service.

6.2 Rendezvous Protocol

6.2.1 Introduction

The Rendezvous Protocol (RVP) is used for propagation of messages within a peer group. The Rendezvous Protocol provides mechanisms which enable propagation of messages to be performed in a controlled way. To most efficiently propagate messages some peers may agree to do extra work. Each *Rendezvous Peer* cooperates with other Rendezvous Peers and with client peers to propagate messages amongst the peers of a peer group. The Rendezvous Peers work together to form a PeerView. The PeerView is a list of the peers which are currently acting as Rendezvous Peers. The PeerView is structured such that Rendezvous Peers are able to direct messages within the peer group in a consistent way without the need for centralized coordination.

The Rendezvous Protocol is divided into three parts; a management protocol (PeerView), a connection registration protocol and the protocol used for transmitting propagated messages.

The PeerView protocol is an optional protocol used by Rendezvous Peers to organize themselves. The PeerView protocol is used to form and verify the integrity of the PeerView. Each Rendezvous Peer maintains its own local list of active PeerView members. Rendezvous Peers exchange PeerView messages with the other PeerView members to maintain their local PeerView.

The Rendezvous Lease Protocol is an optional protocol which enables non-Rendezvous Peers to subscribe to receive propagated messages. Rendezvous Peers offer leases to subscribing peers based upon the Rendezvous Peer's ability to distribute propagated messages. During the term of the offered lease the Rendezvous Peer will accept messages from the subscribing peer for propagation to the peer group and will distribute any appropriate propagation messages to the subscriber. Both the Rendezvous Peer and the subscriber peer may terminate the lease at any time.

The Rendezvous Propagation Protocol is the only required protocol for peers which use the Rendezvous Service. The Rendezvous Propagation Protocol enables peers to manage the propagation of individual messages within a peer group. The Rendezvous Propagation Protocol enables peers to determine the source of a propagated message, the peers the message has already traversed, the listener name to which the message is addressed and the remaining distance which the message may travel before expiring.

6.2.2 Rendezvous Advertisement

A Rendezvous advertisement describes a peer that acts as a Rendezvous Peer for a given PeerGroup. Those advertisements can be published and retrieved, so peers that are looking for rendezvous peers can find them.

```
<xs:element name="RdvAdvertisement" type="jxta:RdvAdvertisement"/>

<xs:complexType name="RdvAdvertisement">
  <xs:sequence>
    <xs:element name="RdvGroupId" type="jxta:JXTAID" />
    <xs:element name="RdvPeerId" type="jxta:JXTAID" />
    <xs:element name="RdvServiceName" type="xs:string" />
    <xs:element name="Name" type="xs:string" />
    <xs:element name="RdvRoute" type="jxta:RA" minOccurs="0" />
  </xs:sequence>
</xs:complexType>
```

Figure 6.3: Rendezvous Advertisement Schema

<RdvGroupId>

This is a required element that contains the ID of the PeerGroup for which the peer is a rendezvous.

<RdvPeerId>

This is a required element that contains the ID of the peer which is a rendezvous.

<RdvServiceName>

This is a required element that identifies the PeerView to which the Rendezvous Peer belongs.

<Name>

This is an optional name associated with the rendezvous peer. Often the same as the peer name.

<RdvRoute>

An optional route to the Rendezvous Peer. A [Route Advertisement](#) is included as a child of the <RdvRoute> element.

6.2.3 PeerView Protocol

The PeerView Protocol is used by Rendezvous Peers to organize orderly interactions amongst themselves. The PeerView is a collection of the participant peers. The collection of all Rendezvous Peers within a peer group is the Rendezvous PeerView. The state of the peer group's PeerView is not maintained on a specific peer. Instead, each Rendezvous Peer maintains a local version of its view of the global PeerView. Each local version is loosely consistent with all other local versions. The goal of the PeerView Protocol is to create and maintain a consistent Peerview state.

The PeerView participants exchange messages with other PeerView participants. Each PeerView message contains a [Rendezvous Advertisement](#) and is either a probe or response message. Probe messages contain the Rendezvous Advertisement for the peer which initiates the probe. PeerView response messages may contain the Rendezvous Advertisement of the responding peer or Rendezvous Advertisement of another peer. Responses containing advertisements for other than the responding peer is how Rendezvous Peers learn of other PeerView members.

Edge Peer Flag Element (optional)	
Message Element Namespace	jxta
Message Element Name	PeerView.EdgePeer
Mime Media Type	text/plain; charset=UTF-8
Element Content	true

Cached Flag Element (optional)	
Message Element Namespace	jxta
Message Element Name	PeerView.Cached
Mime Media Type	text/plain; chars
Element Content	true

Failure Notification Flag Element (optional)	
Message Element Namespace	jxta
Message Element Name	PeerView.Failure
Mime Media Type	text/plain; chars
Element Content	true

Probe Element (optional)	
Message Element Namespace	jxta
Message Element Name	PeerView.PeerAdv
Mime Media Type	text/xml; charset
Element Content	RdvAdvertisement of the

Response Element (optional)	
Message Element Namespace	jxta
Message Element Name	PeerView.PeerAdv.Response
Mime Media Type	text/xml; charset=UTF-8
Element Content	RdvAdvertisement of the Responding Peer.

Figure 6.4: PeerView Message

The same basic PeerView Message is used for all PeerView interactions. The core of the PeerView Message is either the probe or response element. Probe messages normally result in a response message from the receiving peer. Response messages normally

receive no response. The PeerView Message elements other than the probe or response serve as annotations upon the basic message.

Edge Probe	
Purpose	A probe message by a peer which is not a participant in the PeerView to learn of PeerView participants. The Rendezvous Advertisement included is provided to enable the receiver to respond. It is not actually an advertisement for a Rendezvous Peer. The response should be a PeerView Message containing the Rendezvous Advertisement of random PeerView participant.
Elements	EDGE, PROBE
Sent via	Pipe

Participant Probe	
Purpose	A probe message by a participant peer to another PeerView participant. Response should be a PeerView Message containing the Rendezvous Advertisement of the responding peer. Additionally, a second response containing the Rendezvous Advertisement of a random PeerView participant may also be sent to the probing peer.
Elements	PROBE
Sent via	Endpoint Listener

Participant Response	
Purpose	A response message by a participant peer to another PeerView participant. The response Rendezvous Advertisement must be that of the responding peer.
Elements	RESPONSE

Participant Referral Response	
Purpose	A response message by one PeerView participant peer to another participant peer providing a referral to another different participant peer. The response Rendezvous Advertisement must be that a participant peer which is neither the responding peer nor the intended recipient peer.
Elements	CACHED, RESPONSE
Sent via	Endpoint Listener

Shutdown Notification	
Purpose	A response message by a participant peer providing notification that the responding peer is shutting down and will no longer be available. This is a courtesy message which enables the other participant peers to more quickly update their local views.
Elements	FAILURE, RESPONSE
Sent via	Pipe

Failure Notification	
Purpose	A response message by a participant peer providing notification that the identified peer has failed and is no longer be available. This is a courtesy message which enables the other participant peers to more quickly update their local views.
Elements	CACHED, FAILURE, RESPONSE
Sent via	Pipe

Figure 6.5: PeerView Message Scenarios

6.2.4 Rendezvous Lease Protocol

The Rendezvous Lease Protocol is an optional that allows non-Rendezvous Peers to subscribe to receive propagated messages. The subscriber peers requests a lease of a Rendezvous Peer and may receive a lease in response. Leases are of fixed duration. During the term of the lease the Rendezvous Peer will accept messages from the subscriber peer for propagation to the peer group and will send to the subscriber peer propagated messages received from the peer group.

6.2.4.1 Lease Request Message

When a peer wants to connect to a Rendezvous Peer, it sends a message with the a message element named `jxta:Connect` which contains its Peer advertisement.

6.2.4.2 Lease Granted Message

When a rendezvous peer grants a lease (upon a lease request), it sends a message to the source of the lease request, containing the following message elements:

`jxta:ConnectedLease`

This message element contains (in a String representation) the time in milliseconds the lease is granted for. This message element is mandatory.

`jxta:ConnectedPeer`

This message element contains the PeerID of the rendezvous peer that has granted the lease. This message element is mandatory.

`jxta:RdvAdvReply`

This message element contains the Peer Advertisement of the rendezvous peer that grants the lease. This message element is optional.

6.2.4.3 Lease Cancel Message

When a peer wants to cancel a lease, it sends a message with the following message element:

- "`jxta:Disconnect`": This message element contains the Peer Advertisement of the peer which is requesting to cancel the lease. This message element is mandatory.

6.2.5 Message Propagation Protocol

6.2.6 Behaviour

6.2.6.1 Peer connection

Rendezvous peers which re-propagate messages they have received. A peer can dynamically become a rendezvous peer or may connect to a rendezvous peer. The connection between a peer and a rendezvous peer is achieved by an explicit connection, associated to a lease.

This connection is performed by sending messages using the JXTA Endpoint Protocol. Each peer implementing the Rendezvous protocol must register with the Endpoint Service to receive messages with the following Service Name and Service Param:

- service name: `JxtaPropagate`
- service param: PeerGroup ID

A set of queries and responses are defined by the Rendezvous Protocol in order to establish connections:

- *LeaseRequest* This request is sent by a peer that desires to connect to a given rendezvous. A lease may always be canceled by either party at anytime if necessary. A rendezvous that grants a lease returns *LeaseGranted*.
- *LeaseGranted* This message is sent by a rendezvous to grant a lease to a given client. The amount of time the lease is granted for is included in the message.

6.2.6.2 Propagation Control

The Rendezvous Protocol implementation is responsible for controlling the propagation of messages. A Rendezvous Service should propagate a message unless of the following conditions is detected:

- Loop: if a propagated messages has already been processed on a peer, it is discarded.
- TTL: propagated messages are associated with a Time To Live (TTL). Each time a propagated message is received on a peer, its TTL is decreased by one. When the TTL of a message drops to zero, the message is discarded.
- Duplicate: each propagated message is associated with a unique identifier. When a propagated message has been duplicated, and has already been received on a peer, duplicates are discarded.

This control is performed by embedding a Message Element within each propagated message that is defined as:

```
<xs:element name="RendezVousPropagateMessage" type="jxta:RendezVousPropagateMessage"/>
<xs:complexType name="RendezVousPropagateMessage">
  <xs:sequence>
    <xs:element name="MessageId" type="xs:string" />
    <!-- This should be a constrained subtype -->
    <xs:element name="DestSName" type="xs:string" />
    <xs:element name="DestSParam" type="xs:string" />
    <xs:element name="TTL" type="xs:unsignedInt" />
    <xs:element name="Path" type="jxta:JXTAID" maxOccurs="unbounded" />
  </xs:sequence>
</xs:complexType>
```

Figure 6.6: RendezVous Propagate Message Schema

6.3 Peer Information Protocol

Once a peer is located, its capabilities and status may be queried. PIP provides a set of messages to obtain a peer status information. PIP is an optional JXTA protocol. Peers are not required to respond to PIP requests.

A reliable transport is optional for PIP. Multiple peer information messages may be sent. None, one or multiple responses may be received in response to any query.

The PIP is layered upon the [Peer Resolver Protocol](#). The `<QueryID>` element is used to match PIP queries containing `<request>` elements to the PIP Response Messages containing the matching responses.

6.3.1 Obtaining PIP Responses

The PIP Query Message provides a *request* field that may be used to encode a specific request. PIP does not dictate the format of the *request* field and it is left up to the consumer to do so. Higher-level services may utilize the request field to offer expanded capabilities.

6.3.2 PIP Query Message

The query message is sent to a peer to query the current state of the peer, and obtain other relevant information about the peer. A query without a defined request field returns a default set of information about a peer (i.e. uptime, message count, etc.).

```
<xs:element name="PeerInfoQueryMessage" type="jxta:PeerInfoQueryMessage"/>

<xs:complexType name="PeerInfoQueryMessage">
  <xs:sequence>
    <xs:element name="sourcePid" type="jxta:JXTAID" />
    <xs:element name="targetPid" type="jxta:JXTAID" />
    <!-- if not present then the response is the general peerinfo -->
    <xs:element name="request" type="xs:anyType" minOccurs="0" />
  </xs:sequence>
</xs:complexType>
```

Figure 6.7: PIP Query Message

<sourcePid>

The peer id of the requesting peer.

<targetPid>

The peer id of the peer being queried.

<request>

An optional Request structure.

6.3.3 PIP Response Message

The Peer Information Protocol Response Message provides specific information about the current state of a peer, such as uptime, inbound and outbound message count, time last message received, and time last message sent.

```

<xs:element name="PeerInfoResponseMessage" type="jxta:PeerInfoResponseMessage"/>

<xs:complexType name="PeerInfoResponseMessage">
  <xs:sequence>
    <xs:element name="sourcePid" type="jxta:JXTAID" />
    <xs:element name="targetPid" type="jxta:JXTAID" />
    <xs:element name="uptime" type="xs:unsignedLong" minOccurs="0" />
    <xs:element name="timestamp" type="xs:unsignedLong" minOccurs="0" />
    <xs:element name="response" type="xs:anyType" minOccurs="0" />
    <xs:element name="traffic" type="jxta:piptraffic" minOccurs="0" />
  </xs:sequence>
</xs:complexType>

<xs:complexType name="piptraffic">
  <xs:sequence>
    <xs:element name="lastIncomingMessageAt" type="xs:unsignedLong" minOccurs="0" />
    <xs:element name="lastOutgoingMessageAt" type="xs:unsignedLong" minOccurs="0" />
    <xs:element name="in" type="jxta:piptrafficinfo" minOccurs="0" />
    <xs:element name="out" type="jxta:piptrafficinfo" minOccurs="0" />
  </xs:sequence>
</xs:complexType>

<xs:complexType name="piptrafficinfo">
  <xs:sequence>
    <xs:element name="transport" maxOccurs="unbounded">
      <xs:complexType>
        <xs:simpleContent>
          <xs:extension base="xs:unsignedLong">
            <xs:attribute name="Expiration" type="xs:anyURI" />
          </xs:extension>
        </xs:simpleContent>
      </xs:complexType>
    </xs:element>
  </xs:sequence>
</xs:complexType>

```

Figure 6.8: PIP Response Message

<sourcePid>

The peer id of the requesting peer.

<targetPid>

The peer id of the peer being queried.

<uptime>

The relative time in milliseconds since the responding Peer Information Service began execution. Peers should provide this tag in all responses, but may chose to not implement it if the information is unavailable or would represent a security breach.

<timestamp>

The absolute time at which this response was generated. Measured in milliseconds since "the epoch", namely January 1, 1970, 00:00:00 GMT. Peers should provide this tag in all responses, but may chose to not implement it if the information is unavailable or would represent a security breach.

<response>

Potentially contains a response to a previous request from a PIP Query. To match queries to responses the *QueryId* element of the [Peer Resolver Protocol](#) must match. This field can contain any desired content.

<traffic>

Contains information about the network traffic performed by the target peer. This element is optional.

<lastIncomingMessageAt>

The absolute time at which this peer last received a valid JXTA message on one of its transports. Measured in milliseconds since "the epoch", namely January 1, 1970, 00:00:00 GMT. Peers should provide this tag in all responses, but may chose to not implement it if the information is unavailable or would represent a security breach.

<lastOutgoingMessageAt>

The absolute time at which this peer last sent a valid JXTA message on one of its transports. Measured in milliseconds since "the epoch", namely January 1, 1970, 00:00:00 GMT. Peers should provide this tag in all responses, but may chose to not implement it if the information is unavailable or would represent a security breach.

<in>

If present, contains elements which describe incoming traffic from various endpoint addresses.

<transport>

Provides the number of bytes received by the named endpoint address.

<out>

If present, contains elements which describe outgoing traffic from various endpoint addresses.

<transport>

Provides the number of bytes transmitted by the named endpoint address.

6.4 Pipe Binding Protocol

The Pipe Binding Protocol (PBP) is used by applications and services in order to communicate with other peers. A pipe is a virtual channel between two endpoints described in a Pipe Advertisement. There are two ends of a Pipe: the Input Pipe (receiving end) and the Output Pipe (sending end).

The Pipe Binding Protocol is layered upon the Endpoint Protocol, and will use a variety of Message Transports such as the JXTA HTTP Transport, the JXTA TCP/IP Transport, or the secure JXTA TLS Transport for the sending of messages.

A pipe can be viewed as an abstract named message queue, supporting create, open/resolve (bind), close (unbind), delete, send, and receive operations. Actual pipe implementations may differ, but all compliant implementations use the PBP to bind the pipe to an endpoint.

A reliable message transport is optional. Multiple binding query messages may be sent. None, one or multiple responses may be received.

6.4.1 Pipe Advertisement

A Pipe Advertisement describes a pipe. The pipe advertisement is used by the pipe service to create associated local input and output pipe endpoints.

Each pipe advertisement includes a Pipe ID which is the canonical name for the pipe.

Each pipe advertisement must include a pipe type. There are currently three different types of pipes:

- **JxtaUnicast** Unicast, unsecure and unreliable. This type of pipe is used to send one-to-one messages.
- **JxtaUnicastSecure** Unicast, secure (using TLS). An extension of the **JxtaUnicast**, except that the data is protected using a virtual TLS connection between the endpoints.
- **JxtaPropagate** Diffusion pipes. This pipe type is used to send one-to-many messages. Any peer that has enabled an Input Pipe on a propagate type pipe receives messages that are sent onto it.

A pipe advertisement may include an optional symbolic name.

```

<xs:element name="PipeAdvertisement" type="jxta:PipeAdvertisement"/>

<xs:complexType name="PipeAdvertisement">
  <xs:sequence>
    <xs:element name="Id" type="jxta:JXTAID" />
    <xs:element name="Type" type="xs:string" />
    <xs:element name="Name" type="xs:string" minOccurs="0" />
    <xs:element name="Desc" type="xs:anyType" minOccurs="0" />
  </xs:sequence>
</xs:complexType>

```

Figure 6.9: Pipe Advertisement Schema

<Id>

This is a required element that uniquely identifies the pipe. Each pipe has a unique id. See [IDs](#) for description of JXTA Ids.

<Type>

This is an required element that defines the type of the pipe. The following types are currently defined:

JxtaUnicast

may not arrive at the destination, may be delivered more than once to the same destination, may arrive in different order.

JxtaUnicastSecure

may not arrive at the destination, may be delivered more than once to the same destination, may arrive in different order, but is encrypted using TLS.

JxtaPropagate

one to many pipe.

<Name>

This is an optional name that can be associated with a pipe. The name is not required to be unique unless the name is obtained from a centralized naming service that guarantee name uniqueness.

Example 6.3 Pipe Advertisement

```

<?xml version="1.0" encoding="UTF-8"?>

<!DOCTYPE jxta:PipeAdvertisement>

<jxta:PipeAdvertisement xmlns:jxta="http://jxta.org">
  <Id>urn:jxta:uuid-094AB61B99C14AB694D5BFD56C66E512FF7980EA1E6F4C238A26BB362B34D1F104</Id>
  <Type>JxtaUnicast</Type>
  <Name>Talk to Me!</Name>
</jxta:PipeAdvertisement>

```

6.4.2 Pipe Resolver Message

For some pipe types, notably `JxtaUnicast` and `JxtaUnicastSecure` it is necessary to locate a peer which is listening on the pipe in order to create an Output Pipe. For these pipe types the Pipe Service uses the "Pipe Resolver Message". The same message schema is used for both the resolve query and for the response.

```

<xs:element name="PipeResolver" type="jxta:PipeResolver"/>

<xs:simpleType name="PipeResolverMsgType">
  <xs:restriction base="xs:string">
    <!-- QUERY -->
    <xs:enumeration value="Query"/>
    <!-- ANSWER -->
    <xs:enumeration value="Answer"/>
  </xs:restriction>
</xs:simpleType>

<xs:complexType name="PipeResolver">
  <xs:sequence>
    <xs:element name="MsgType" type="jxta:PipeResolverMsgType"/>
    <xs:element name="PipeId" type="jxta:JXTAID"/>
    <xs:element name="Type" type="xs:string"/>

    <!-- used in the query -->
    <xs:element name="Cached" minOccurs="0" default="true" type="xs:boolean"/>
    <xs:element name="Peer" minOccurs="0" maxOccurs="unbounded" type="jxta:JXTAID" />

    <!-- used in the answer -->
    <xs:element name="Found" minOccurs="0" type="xs:boolean"/>
    <!-- This should refer to a peer adv, but is instead a whole doc -->
    <xs:element name="PeerAdv" minOccurs="0" type="xs:string"/>
  </xs:sequence>
</xs:complexType>

```

Figure 6.10: Pipe Resolver Message Schema

<MsgType>

Used to indicate if it is the Query or the Answer. May be one of:

Query

This is a query.

Answer

This is a response.

<PipeId>

The Pipe ID which is being resolved.

<Type>

The type of pipe resolution being requested. This value must match the value of <Type> from the Pipe Advertisement.

<Cached>

If *false*, peers which do not have the pipe bound locally as an Input Pipe must not respond to the query. They may forward the query to peers which they believe to have the pipe bound as an Input Pipe. This feature is deprecated and implementations should treat the tag as always being *false*.

<Peer>

A peer id. In Queries, if present, it specifies the Peer ID of the only peer from which responses will be expected. Responses from all other peers may be ignored. This does not guarantee a response to the pipe binding request will be made by the peer. Response to pipe binding requests is always optional. In Answer messages, all of the peers on which the Input Pipe is known to be bound.

<Found>

Used to indicate if the Input Pipe was found on the specified peer.

<PeerAdv>

Peer Advertisement of the peer which resolved the Input Pipe. This peer may appear in the list of peer ids on which the Input Pipe is bound, but this should not be assumed.

6.4.3 Propagate Pipe Message Header

Every message sent on a propagate pipe includes a message element which is used to manage the propagation of messages. The message element is stored in the `jxta` namespace and has the name `JxtaWireHeader`.

```
<xs:element name="JxtaWire" type="jxta:JxtaWire"/>

<xs:complexType name="jxta:JxtaWire">
  <xs:sequence>
    <xs:element name="SrcPeer" minOccurs="0" type="jxta:JXTAID" />
    <xs:element name="PipeId" type="jxta:JXTAID" />
    <xs:element name="MsgId" type="xs:string" />
    <xs:element name="TTL" type="xs:unsignedInt" />
    <xs:element name="VisitedPeer" type="jxta:JXTAID" maxOccurs="unbounded" />
  </xs:sequence>
</xs:complexType>
```

Figure 6.11: Propagate Pipe Message Header Schema

<SrcPeer>

The peer which originated this message.

<PipeId>

The ID of the pipe on which the message is being sent.

<MsgId>

Each message has an associated token which is used for duplicate tracking. This token is normally generated using a pseudo random source or partially pseudo randomly in order to reduce the chance of collisions.

<TTL>

Each peer attempting to forward this message should first decrement the TTL value. If the value reaches zero then the message should not be forwarded to any additional peers.

<VisitedPeer>

The set of peers which are known to have seen this message. Peers which forward propagate pipe messages should add themselves to this list before forwarding the message. They should also avoid forwarding the message to any of the peers listed.

Chapter 7

Standard JXTA Message Transports

JXTA defines several Message Transport Bindings for implementing JXTA messaging on top of existing networks.

7.1 TCP/IP Message Transport

7.1.1 Introduction

This section describes the TCP/IP Message Transport, the most commonly used Message Transport by JXTA. It is named the TCP/IP binding because it uses no other underlying protocols whereas other Message Transports build upon, for example, the HTTP "[RFC2616](#)", TLS "[RFC2246](#)", BEEP "[RFC3080](#)" "[RFC3081](#)" and SMTP "[RFC821](#)" protocols. The TCP/IP transport is designed to be the simplest and quickest to implement of the JXTA Message Transport bindings. Unfortunately, the TCP/IP Message Transport can be adversely affected by the partitioning of the Internet Address Space caused by Network Address Translation (NAT), Firewalls, and conflicting use of Private IP Address ranges. For these circumstances one of the other Message Transports may be more efficient.

7.1.2 Choreography

This section describes the protocol exchange implemented by the JXTA TCP/IP Message Transport. The two participants operate symmetrically, the message sequence is the same for both sides.

- Connection Opens
- Welcome Message
- Message(s)
- Connection Closes

The connection may be closed by either peer following the transmission of any number of messages. If a peer discover an error in the transmission (unexpected input, unsupported message content, framing problems, etc.) it must close the connection.

7.1.3 Welcome Message

The Welcome Message is sent by both peers of a JXTA TCP/IP Communication. It's primary purpose is to identify the endpoint address that each peer has associated with TCP/IP connections on that interface/port. Additionally, the welcome message provides a reasonable 'ping' response if the only goal of the connection is to determine connectivity.

A peer should send its Welcome Message as soon as the connection is open, but must not send any Message Bodies until it has received a Welcome Message from the remote peer. A peer must send its peer id as the connection endpoint address. The welcome message can be no longer than 4096 octets in length, though in practice it will almost always be much shorter.

```

<WELCOME>      ::= <GREETING> <SPACE> <WELCOMEDEST> <SPACE>
                  <WELCOMESRC> <SPACE> <WELCOMEPER> <SPACE>
                  <NOPROP> <SPACE> <VERSION> <CRLF>

<GREETING>     ::= "JXTAHELLO"

<SPACE>        ::= octet #x20

<WELCOMEDEST>  ::= <ENDPOINTADDRESS>

<WELCOMESRC>   ::= <ENDPOINTADDRESS>

<WELCOMEPER>   ::= <JXTAID>

<NOPROP>       ::= "0" | "1"

<VERSION>      ::= "1.1"

<CRLF>         ::= octet #x0D octet #x0A

```

Figure 7.1: Welcome Message ABNF

This message is designed to be human readable. This primarily for the purpose of debugging. This message also allows a peer to be ‘pinged’ using Telnet or a web browser.

<WELCOMEDEST>

The assumed destination of messages which will be sent on this connection. If a received message has no destination address then this is the address that should be used.

<WELCOMESRC>

The public address of the JXTA Message Transport instance which is handling this connection.

<WELCOMEPER>

Provides the peer id of the peer you have reached with this connection.

<NOPROP>

If 1 then the remote peer does not wish to receive propagate/broadcast messages on this connection.

Example 7.1 Welcome Message

```
JXTAHELLO tcp://69.3.88.186:34368 tcp://209.25.154.236:9701 urn:jxta:uuid-59616261646162614 ↔
A7874615032503345A8391EC0914B24B264AF31F297A6FD03 1 1.1
```

7.1.4 Message Transmission

Following the [Welcome Message](#), the connection is used to send JXTA Messages. Each message is preceded by a small amount of header information. The header information is similar to HTTP headers, but is binary encoded to provide for lighter-weight parsing.

```

<PACKAGE>      ::= <HEADERBLOCK> <BODY>

<HEADERBLOCK>  ::= (<HEADER>)* <EMPTYHEADER>

<HEADER>       ::= <HEADERNAME> <HEADERBODY>

<HEADERNAME>   ::= octet [length] (octets) [length]

<HEADERBODY>   ::= octet [length msb] octet [length lsb] (octets) [length]

<EMPTYHEADER>  ::= octet #x00

<BODY>         ::= octets

```

Figure 7.2: JXTA Message Package ABNF

The header block will contain both required headers and may contain additional optional headers. The headers are loosely ordered. Headers of a particular type are considered by ordered only with respect to other headers of the same type, they have no order relationship to headers of other types. In general, headers may appear in any order. The header names are case-insensitive canonical UTF8 (NFC) (see *Unicode Standard Annex #15 : Unicode Normalization Forms "USA15"*) strings.

7.1.4.1 JXTA Message Package - Required Headers

These headers must appear in every message transmission and have no default values.

content-length

the length of the <BODY> portion. If `content-coding` headers are included, this length is still the transfer length, not the logical length of the body. The header body is an eight octet 'network byte order' length.

content-type

The logical type of the message body expressed as a MIME-type. The header body consists of a MIME type encoded in canonical UTF8 (NFC) using the presentation and encoding of *Multipurpose Internet Mail Extensions (MIME) Part Two: Media Types "RFC2046"*, ie. any encoding specified by RFC 2046 is performed before the string is encoded into UTF8 from the implementation native string representation.

Peers must treat unrecognized `content-type` specifications as an error and close the connection.

7.1.4.2 JXTA Message Package - Optional Headers

These headers may appear in any message transmission and have no default values.

content-coding

Specifies a coding which has been applied to the logical message body prior to transmission. The header body contains an eight octet 'network byte order' value which was the length of the message body before encoding. This is followed by the Mime-type of the content coding.

`content-coding` headers are listed in the order in which they were applied, ie. The first `content-coding` header was the first coding applied and will be the last to be removed.

All other possible headers are also optional. Unrecognized optional headers should be ignored.

7.1.5 IP Multicast Usage

The TCP Transport may be configured to accept broadcast messages through IP Multicast. Messages received are not part of an ongoing connection between peers. Each broadcast message is idempotent. The datagram is formed similarly to that used for unicasts connections. It consists of :

- A Welcome Header
- A Message Package Header
- A single message

Additionally, in order to fully specify the source of the message, the Message Package Header must contain the following header field.

srcEA

The public addresss of the transport which is sending this broadcast message.

7.2 HTTP Message Transport

The JXTA HTTP Message Transport is logically divided into two functions: the initiator and the receiver. The initiator is able to establish connections, whereas the receiver is able accept connections. A peer may provide either function or both.

The HTTP Message Transport is designed to work effeciently with or without HTTP 1.1 persistent connections and also impose a minimum of overhead beyond that required by HTTP itself.

7.2.1 The HTTP Initiator

The HTTP Initiator begins the connection between two peers. The connection is first used to determine the logical identities of the participating peers. In most cases, this is simply a verification as the connection address was chosen by the initiator to communicate with the desired peer.

The HTTP Initiator uses two types of HTTP request to communicate with the recevier. `GET` requests are used for determining the logical identity of the HTTP Receiver and for polling the HTTP Receiver for messages. `POST` requests are used for sending and receiving JXTA messages with the receiver.

7.2.2 The HTTP Receiver

The HTTP Receiver accepts incoming connections from HTTP initiators. The HTTP Receiver then acts in response to HTTP requests from the initiator.

The HTTP Receiver accepts two types of HTTP request from the initiator, `GET` requests and `POST` requests. `GET` requests can be used for two tasks: determining the logical identity of the receiver, polling for messages destined to the initiator. `POST` requests are used to send messages to the receiver and optionally to receive messages from the receiver.

7.2.3 HTTP Messages

7.2.3.1 'Ping' Command

The 'ping' command is used to determine the logical identity (peer id) of a receiver. It also has the effect of determining connectivity (hence the name).

The HTTP initiator issues a request as follows:

Request

`GET`

Resource

`/`

The HTTP Receiver should respond to well-formed requests as follows:

Response Code

200

Required Headers**Content-Length**

Must always be present.

Content-TypeMust be `text/plain; charset="UTF-8"`.**Response Body**

The peer id of the HTTP receiver expressed as a JXTA Endpoint Address.

Example 7.2 HTTP Ping Example

Request:

GET / HTTP/1.1

Response:

200

content-length=61

content-type=text/plain; charset="UTF-8"

jxta://urn:jxta:uuid-59616261646162614A7874615032503304BD268FA4764960AB93A53D7F15044503

7.2.3.2 'Poll' Command

The 'poll' command is used by the HTTP initiator to retrieve a message from the HTTP receivers.

The HTTP initiator issues a request as follows:

Request

GET

Resource`/<initiator peerid>?<responseWait>,<extraResponseWait>,<destAddr>`**initiator peerid**The unique portion of the peerid of the initiator, ie. without the `urn:jxta:` prefix.**responseWait**

The maximum duration in milliseconds which the HTTP Receiver should wait in order to send a response. Expressed as a decimal integer. This is the longest duration that the initiator is willing to wait for a response. If unspecified, the HTTP receiver must assume a value of 0 (zero). A delay value of 0 (zero) implies that the HTTP Receiver should wait indefinitely until a response is available. Negative wait values imply that the receiver must not return any response.

lazyClose

The maximum duration in milliseconds during which the HTTP Receiver should send additional response messages. Expressed as a decimal integer. If unspecified the receiver must assume a value of 0 (zero). A delay value of 0 (zero) implies that the HTTP receiver should continue to return message responses until the connection is closed. Negative values imply that the receiver must not return any additional responses.

destAddr

A URI specifying the complete destination address for this connection. If there is no destination specified, the receiver must assume a value of the address of the receiving network interface and receiving port. The destination service and service parameter are both uninitialized.

The HTTP Receiver should respond to well-formed requests as follows:

Response Code

200

Required Headers**Content-Type**

Must match the content type of the JXTA Message Body which forms the response, if any.

Response Body

Well-formed JXTA Message Bodies

Example 7.3 HTTP Poll Example

Request:

```
POST /uuid-59616261646162614A7874615032503304BD268FA4764960AB93A53D7F15044503?0,0,http ↔
: //64.81.53.91:8720/EndpointService:jxta-NetGroup/uuid- ↔
DEADBEEFDEAFBABA FEEDBABE0000000F05/connect,29988,keep,true HTTP/1.1
```

Response:

```
200
content-length=0
```

7.2.3.3 'Send' Command

The 'send' command is used by the HTTP initiator to send a message to the HTTP receiver and optional retrieve a response message.

The HTTP initiator issues a request as follows:

Request

POST

Resource

/<initiator peerid>

initiator peerid

The unique portion of the peerid of the initiator, ie. without the `urn:jxta:` prefix.

Required Headers**Content-Length**

Must always be present.

Content-Type

Must match the content type of the JXTA Message Body which forms the response.

Request Body

A well-formed JXTA Message Body

The HTTP Receiver should respond to well-formed requests as follows:

Response Code200

Example 7.4 HTTP Send Example

Request :

```
POST /uuid-59616261646162614A7874615032503304BD268FA4764960AB93A53D7F15044503 HTTP/1.1
```

Response :

```
200
```

7.3 TLS Transport Binding

7.3.1 Introduction

JXTA implements *IETF RFC 2246 : Transport Layer Security Version 1 "RFC2246"* on top of the [Endpoint Router Transport Protocol](#) using JXTA messages. Messages are transmitted across one of the other Transport Bindings. This implementation provides for secure message delivery between Peer Endpoints even when multiple hops across JXTA peers must be used.

7.3.2 TLS Messages

The TLS protocol exchange is accomplished between the participant peers using JXTA Messages. Three JXTA Message Element types are used to implement the protocol:

- TLS Contents
- Acknowledgements
- Retries

7.3.2.1 TLS Content Element

Used to transfer TLS content between peers.

Element Name

UTF-8 String containing the sequence number of this element. Sequence numbers are monotonically increasing positive integers.

Element Type

application/x-jxta-tls-block

Element Data

TLS Protocol Data per *IETF RFC 2246 "RFC2246"*

7.3.2.2 TLS ACK Element

Used to transfer TLS content between peers.

Element Name

TLSACK as a UTF-8 String.

Element Type

application/x-jxta-tls-ack

Element Data

Sequential ACK

Network Byte Order Unsigned 32 Bit Integer. Indicates that the remote peer has successfully received all of the messages with sequence numbers up to and including this value.

Selective ACKs (optional)

Network Byte Order Unsigned 32 Bit Integer(s). Each entry is the sequence number of a message which has been successfully received by the remote peer. Entries are all greater than the sequential ACK value and in numerically increasing order.

7.3.2.3 Retry Element

Used to signify that this message is being sent as a retry of an apparently failed previous send attempt. If present, this element will accompany a [TLS Content Element](#). The element is mostly informational, but can be used to assist in window control.

Element Name

MARKRetr as a UTF-8 String.

Element Type

text/plain; charset="UTF-8"

Element Data

TLSRET as a UTF-8 string.

Chapter 8

JXTA Message Wire Representations

In order to send JXTA messages between peers JXTA a serialized ‘wire’ encoding of the message must be created. The standard JXTA *Message Transport* implementations use a common set of wire encodings. There are two standard message encoding formats for JXTA Messages, XML and binary. Each JXTA Message Transport uses the encoding format most appropriate for its’ capabilities.

8.1 General Requirements

Message encoding formats for Message Transports must allow for the faithful transmission of arbitrary messages.

Message Elements with duplicate names must be supported.

8.2 Binary Message Format

XML and other textual representations such as MIME ["RFC2045"](#) are inefficient for the transmission of arbitrary application data. The JXTA Binary Message Format is designed to facilitate the efficient transmission of data between peers.

Message Transports which are capable of sending and receiving unencoded binary data will probably prefer to use the Binary Message Format. This format is designed such that all components are of declared size, no parsing is ever required to determine lengths.

8.2.1 Conventions

Multi-byte lengths are sent with the high order byte first (also known as "Big Endian" or "Network Byte Order"). All strings start with a two byte length, followed by a UTF string value. The message format is specified using ABNF ["RFC2234"](#). ABNF is normally used for ASCII grammars, but here we use it to define a byte sequence for a binary message.

8.2.2 Binary Message Version 1

```

msg      ::= "jxmg"                ; signature (0x6a 0x78 0x6d 0x67)

          version                    ; One byte. "0".

          namespaces

          element_count              ; two bytes (binary)

          1* elm

namespaces ::= namespace_count      ; two bytes (binary)

          0* namespace              ; Each namespace is a string

namespace ::= string

string    ::= len2                  ; two bytes (binary)

          len2 * UTF8 chars          ; characters

elm       ::= "jxel"                ; signature (0x6a 0x78 0x65 0x6c)

          namespaceid               ; one byte (binary)

          flags                      ; Indicates which parts follow (binary)

          name                      ; element name

          [type]                    ; Present if (flags & HAS_TYPE)

          [encoding]                ; Present if (flags & HAS_ENCODING)

          len4                      ; Four byte binary length of content

          content                    ; element content

          [signature]               ; associated signature element
                                   ; Present if (flags & HAS_SIGNATURE)

name      ::= string

type      ::= string                ; Mime Media Type

encoding  ::= string                ; Mime Media Type

content   ::= len4 * byte           ; the bytes of the content.

signature ::= elm

flags     ::= byte                  ; HAS_TYPE      = 0x01;
                                   ; HAS_ENCODING   = 0x02;
                                   ; HAS_SIGNATURE  = 0x04;

```

Figure 8.1: JXTA Binary Message Version 1 ABNF

8.2.2.1 Message Header

Each message starts with the four byte signature `jxmj`. (0x6a 0x78 0x6d 0x67) The signature is used to aid in sanity checking a transport implementation. This is followed by a one byte version number which must 0 (zero). Next is a list of namespaces used by this message. See the production rule for namespaces below. And last is a two byte element count followed by the elements themselves.

8.2.2.2 Namespaces

Each message element a member of a namespace. The namespaces of all elements are placed into an ordered list at the start of the message. Each entry in the list is assigned an id. The first entry in the list is assigned an id of 2. The id of each successive namespace is one plus the id of the preceding namespace. The ids 0, and 1 are preassigned. 0 represents the empty namespace. 1 is the `jxta` namespace. See [Namespace](#) for information regarding use of namespaces.

The namespaces portion of a binary message starts with a two byte namespace count. The count is followed by a sequence of namespace names. It is permissible for this sequence to be empty.

8.2.2.3 Message Element Header

Each message elements starts with the four byte signature `jxel` (0x6a 0x78 0x65 0x6c). The signature is used to aid in sanity checking a transport implementation. Next is the namespace id byte. This byte indicates which name space this element belongs to. The next byte is the flags byte. Bits in this flag byte indicate which of the optional components of an element are present. `HAS_ENCODING` is currently unsupported.

The flags are followed by the element name. The MIME Media Type of the element, if present, follows the element name. If the type is not specified for an element then the type `application/octet-stream` is assumed. The element type is optionally followed by the encoding type. Next is the four byte length of the content, followed by the content data of the element.

8.2.3 Binary Message Version 2

In mid-2006 a new version of the JXTA Binary Message was developed. This new version currently remains optional for all JXTA implementations. The new version was developed in order to reduce the wire overhead of JXTA Binary Messages and to improve support for larger messages, non-UTF8 character encodings and content encoding.

```

msg      ::= "jxmg"           ; signature (0x6a 0x78 0x6d 0x67)
          version             ; One byte. "1".
          flags               ; UTF-16BE String = 0x01
                              ; UTF-32BE String = 0x02
          namestable          ; Common names used in element headers
          element_count       ; The number of message elements or zero
          element_count * element

namestable ::= names_count      ; Number of names defined in the table.
          ; Does not include implicit names.
          names_count * string ; names_count names

names_count ::= 2 * byte        ; 2 byte MSB

string      ::= len2           ; string length
          len2 * chars         ; characters (UTF-8, UTF-16BE, UTF-32BE per flags)

element     ::= "jxel"         ; signature (0x6a 0x78 0x65 0x6c)
          flags               ; Flags for Message Element
                              ; Body lengths are UINT64 = 0x01
                              ; Element Name is literal = 0x02
                              ; Element has MIME Type = 0x04
                              ; Element has Signature Element = 0x08
                              ; Element has Content Encodings = 0x10
                              ; Signature is of encoded element body = 0x20
          namespaceid         ; name
          (name | string)     ; (see flag 0x02)
          [type]              ; (see flag 0x04)
          [encodings]         ; (see flag 0x10)
          elementLen          ; length of content (see flag 0x01)
          content             ; element content
          [signature]         ; associated signature element
                              ; (see flag 0x10)

type        ::= name           ; Mime Media Type

encodings   ::= num_encodings  ; number of encodings
          num_encodings * encoding

encoding    ::= elementLen     ; pre-encoded element length
          name                 ; encoding mime type.

num_encodings ::= len1         ; 1 byte length

elementLen  ::= (len4 | len8)  ; len4 or len8 (see flag 0x01)

content     ::= elementLen * byte ; the bytes of the content.

signature   ::= element        ; signature message element

len1        ::= 1 * byte       ; 1 byte length

len2        ::= 2 * byte       ; 2 byte length (MSB order)

len4        ::= 4 * byte       ; 4 byte length (MSB order)

len8        ::= 8 * byte       ; 8 byte length (MSB order)

name        ::= (2 * byte | 0xFFFF {string})

```

Figure 8.2: JXTA Binary Message Version 2 ABNF

8.2.3.1 Message Header

Each JXTA Binary Message begins with a header that provides basic parameters and options for the message. The message header starts with a four byte signature, `jxmg` (0x6a 0x78 0x6d 0x67). The signature is used to aid in sanity checking a transport implementation. This is followed by a one byte version number which must be 1. The version number is followed by message flags. These flags apply to all message elements included in the message. Following the flags is a list of names used by this message. The Message Header ends with a count of the Message Elements which will follow. The element count may also be zero indicating that the message elements are uncounted.

8.2.3.2 Names Table

Every message includes a table of names (strings). The names table provides a mechanism for looking up commonly used Message Element header strings. Each entry in the list is assigned an ID. The first entry in the list is assigned an id of 2. The ID of each successive name is one plus the id of the preceding names. The IDs 0, 1, and 0xFFFF are preassigned. 0 represents the empty namespace. 1 is the `jxta` namespace. 0xFFFF is used when a literal name must be used within a Message Element. Generally this happens when messages are created on-the-fly and the name was not known at the time the Message was begun.

The names table portion of a JXTA Binary Message starts with a two byte names count. The count is followed by a sequence of name strings. It is permissible for this sequence to be empty.

8.2.3.3 Message Element Header

Each Message Element header starts with the four byte signature `jxel` (0x6a 0x78 0x65 0x6c). The signature is used to aid in sanity checking a transport implementation. The signature is followed by the Message Element flags. The flags control Message Element specific options. The flags are followed by the namespace ID. Each Message Element is part of a specific namespace. The namespace ID identifies the name of the namespace. The namespace is followed by the name of the Message Element. Depending upon the setting of flag #1 this may be either a name ID or a string literal. String literals are commonly used for names which only occur once within the context of a message. Optionally following the Message Element Name is the Message Element Type. The Message Element Type identifies the MIME multimedia type associated with this message element's content. If no Message Element Type is included then the MIME type is assumed to be `application/octet-stream`. Following the type is the optional content encoding information. If present the content encoding information describes any content transformations which have been applied to the Message Element content and must be reversed in order to recover the content. Examples of content encoding include compression, chunking and encryption. Following the content encoding information is the Message Element content length. The content length specifies the size in bytes of the encoded Message Element content. The content length is followed by the Message Element content which may need to be unencoded as it is processed.

8.3 XML Message Format

The XML message format is used for transports which can transmit text but not raw binary data. An effort is made to keep the message elements as readable as possible. See the section on encoding.

8.3.1 Message

Example 8.1 XML Message

```
<?xml version="1.0"?>

<!DOCTYPE Message>

<Message version="0">
  <Element name="jxta:SourceAddress" mime_type="text/plain">
    tcp://123.456.205.212
  </Element>
  <Element name="stuff" encoding="base64" mime_type="application/octet-stream">
    AAECAwQFBgcICQoLDA0ODxAREhMUFRYXGBkaGxwdHh8gISIjJCUmJygpKissLS4vMDEyMzQ1Njc4
    OTo7PD0+P0BBQkNERUZHSElKS0xNTk9QUVJTVFVWV1hZWltdcXV5fYGFfiY2RlZmdoaWprbGlub3Bx
    cnN0dXZ3eHl6e3x9fn+AgYKDhIWGh4iJiouMjY6PkJGSk5SVlpeYmZqbnJ2en6ChoqOkpaangKmq
    q6ytrq+wsbKztLW2t7i5uru8vb6/wMHCw8TFxsc=
  </Element>
</Message>
```

The top level XML element is the Message element. There is one required attribute, version. The value of version must be zero "0". The body of the Message element is a sequence of Element elements.

8.3.2 Element

Each Element must have a name and mime_type attribute. Optionally an encoding attribute may be present.

8.3.2.1 Name

This is a required attribute names the element. The name contains the namespace, followed by a colon, followed by the simple name of the element.

8.3.2.2 MIME type

A required attribute indicating the MIME type of the element. If the MIME type was not specified in the message, a type of "application/octet-stream" is used.

8.3.2.3 Encoding

Encoding is optional. The only supported encoding at this time is base64. *IETF RFC 1521* "[RFC1521](#)". If the encoding is not present, the content is just treated as a string. While name, type and content are parts of a message which will be found in all implementations, the encoding is added by this particular format.

Note: The reference implementation uses the following technique to select encoding. If the MIME major type is "text", the content is treated as a string. Therefore, no encoding attribute is present on the element. Otherwise the content is base64 encoded. This will result in human readable text where possible. The content will successfully be transferred using an XML message even if the user has incorrectly specified the MIME type. There is no requirement that an implementation follow this rule. Implementations will still interoperate whether or not this rule is used.

Part III

JXTA Reference Implementations Information

Chapter 9

Java 2 SE Reference Implementation

9.1 JXTA ID Formats

These ID Formats have been defined for JXTA Language Binding Reference Implementations.

9.1.1 JXTA ID Formats : 'uuid' ID Format

The Java 2SE reference implementation of JXTA provides an implementation of JXTA IDs based upon UUIDs *ISO/IEC 11578:1996* ["ISO11578"](#). This optional ID Format is identified as the `uuid` format. In this implementation, UUIDs are encoded as hex digits as the basis generating unique identifiers. At the end of each UUID ID are two hex characters that identify the type of JXTA ID. Currently, six ID Types have been defined.

```

<JXTAUUIDURN>      ::= "urn:" <JXTANS> ":" <JXTAUUIDFMT> "-"
                      <(1*(<hex> <hex>)) <JXTAUUIDIDTYPE>

<JXTAUUIDFMT>      ::= "uuid"

<JXTAUUIDIDTYPE>   ::= <CODATID> | <PEERGROUPID> | <PEERID> |
                      <PIPEID> | <MODULECLASSID> | <MODULESPECID>

<CODATID>          ::= "01"

<PEERGROUPID>      ::= "02"

<PEERID>           ::= "03"

<PIPEID>           ::= "04"

<MODULECLASSID>    ::= "05"

<MODULESPECID>     ::= "06"

```

Figure 9.1: JXTA "uuid" ID Format ABNF

The characters preceding the ID Type identifier are the encoded form of the ID. The encoding consists of a variable number of characters dependant upon the ID Type being encoded. To decode the ID the hex characters are translated in order and placed into the elements of a byte array from which the various ID components can be retrieved. All of JXTA UUID IDs are currently manipulated as 64 byte arrays although no ID Type currently requires all the full 64 bytes to encode their contents. Position

63 always contains the UUID ID Type value. The remainder of the ID fields are defined beginning at Position 0 and increment towards Position 63.

To make the text presentation of JXTA UUID IDs as URNs more compact implementations must not encode the value of unused Positions within the array. Since they are irrelevant to the value of the ID they can be assumed to be zero. Implementations must also omit from the encoding the value of any Positions that precede the unused portion and contain zero. The reference Java implementation accomplishes this by scanning from Position 62 towards Position 0 searching for the first non-zero value. It then encodes from position 0 to the discovered location followed by the encoding for Position 63. The text encoding of a JXTA ID must be canonical according to the URN specification, thus this ‘zero-saving’ technique must be present in every implementation.

Example 9.1 Sample "uuid" Format ID

```
urn:jxta:uuid-00030102040501
```

Decodes to:

```
0:00  1:03  2:01  3:02  4:04  5:05  6-62:00  63:01
```

9.1.2 JXTA UUID Field Definitions

Each of the six JXTA ID Types has a specific definition for how its fields are represented within the common 64-byte array structure. Common between the six ID Types is the definition of Position 63. This location is reserved for the ID Type.

9.1.2.1 JXTA UUID Codat ID Fields

Each Codat is assigned a unique codat id that enables canonical references to be made to the codat in the context of a specific peer group. A CodatID is formed by the conjunction of a PeerGroupID, a randomly chosen value that has a high probability of being unique, and an optional SHA1 cryptographic hash of the codat contents.

0:Group MSB	0-15:...GROUP UUID...		
...GROUP UUID (cont.)...			15:Group LSB
16:ID MSB	16-31:...CODAT UUID...		
...CODAT UUID (cont.)...			31:ID LSB
32:Hash	32-51:CODAT SHA1 Hash...		
...CODAT SHA1 Hash...			
...CODAT SHA1 Hash		51:Hash	
			63:ID Type

Figure 9.2: JXTA UUID Codat ID Fields

9.1.2.2 JXTA UUID PeerGroup ID Fields

Each peer group is assigned a unique id that enables canonical references to that peer group.

0:MSB	0-15...GROUP UUID...	
...GROUP UUID (cont.)		15:LSB
0:MSB	0-15...PARENT GROUP UUID...	
...PARENT GROUP UUID (cont.)		15:LSB
		63:IDType

Figure 9.3: JXTA UUID PeerGroup ID Fields

9.1.2.3 JXTA UUID Peer ID Fields

Each peer is assigned a unique peer id that enables canonical references to be made to the peer in the context of a specific peer group.

0:Group MSB	0-15...GROUP UUID...	
...GROUP UUID (cont.)		15:Group LSB
16:ID MSB	16-31:...PEER UUID...	
...PEER UUID (cont.)...		31:ID LSB
		63:IDType

Figure 9.4: JXTA UUID Peer ID Fields

9.1.2.4 JXTA UUID Pipe ID Fields

Each pipe is assigned a unique pipe id that enables canonical references to be made to the pipe in the context of a specific peer group.

0	1	2	3	4	5	6	7
0:Group MSB	0-15...GROUP UUID...						
...GROUP UUID (cont.)						15:Group LSB	
16:ID MSB	16-31:...PIPE UUID...						
...PIPE UUID (cont.)...						31:ID LSB	
						63:IDType	
0	1	2	3	4	5	6	7

Figure 9.5: JXTA UUID Pipe ID Fields

valid (the next hop in the list is not reachable), or if the message does not contain a forward route, the local route is then used. If there is no local route, then the message is dropped.

[PENDING TASK] When a router is asked to route a message, and when no route is available, it should search for a route (sending a RouteQuery), and/or send a NACK message back to the source of the message.

When the Endpoint Router receives a RouteQuery, if it knows a route, it answer with a RouteResponse including the forward route.

[PENDING TASK] PingQuery and PingResponse are not yet implemented. In particular, routes should be checked once in a while.

[PENDING TASK] NACK is not yet implemented.

When the Endpoint Router receives an incoming message, and when the incoming message contains a reverse route, the reverse route is added into the local cache of routes. Note that the Endpoint Router detects loops and other errors both in reverse route and forward route.

The Endpoint Router does not remove its message element, even when routed message is eventually delivered locally to the destination application. This allow the application to decide to forward the message to another peer, while keeping the routing information, especially the reverse route, allowing the final destination to respond to the original source of the message without having to issue a RouteQuery.

Chapter 10

Glossary

Advertisement JXTA's language-neutral metadata structures that describe peer resources such as peers, peer groups, pipes, and services. Advertisements are represented as XML documents.

Binding An implementation of the JXTA protocols for a particular environment (e.g., the Java J2SE platform binding).

Codat The combination of a content (commonly a document or file) and a JXTA ID.

Credential A token used to uniquely identify the sender of a message; can be used to provide message authorization.

JXTA JXTA is not an acronym, and in particular the 'J' does not refer to Java. JXTA is a made up word coined by the project's original sponsor, Bill Joy. JXTA is derived from the word Juxtapose, as in side by side. It is a recognition that peer-to-peer is juxtaposed to client server or Web based computing -- what is considered today's traditional computing model.

jux ta pose

tr. v. To place side by side, especially for comparison or contrast.

JXTA Core Specification The JXTA Core Specification consists of the required components and behaviours which are present in all conforming JXTA implementations. This includes the Endpoint Routing Protocol (ERP) and the Peer Resolver Protocol (PRP).

JXTA ID Format A JXTA ID Format is a scheme for representing IDs of JXTA entities. Each ID Format is identified by as sub-namespace of the URN namespace "jxta".

JXTA ID Type A JXTA ID Type describes the characteristics of JXTA IDs which refer to a particular sort of JXTA entity. Currently this includes peer groups, peers, codats, pipes, module classes, module specifications and module implementations, but may be extended to refer to other types of entities in the future or in specific implementations.

JXTA Standard Services The JXTA Standard Services are optional JXTA components and behaviours. Implementations are not required to provide these services, but are strongly recommended to do so. This includes Peer Discovery Protocol (PDP), Peer Information Protocol (PIP), Pipe Binding Protocol (PBP) and Rendezvous Protocol (RVP).

Message Transport A Message Transport is responsible for sending and/or receiving JXTA messages via an external network. A Message Transport may use whatever wire protocol it wishes in order to transmit messages between peers. Message Transports must pass the messages between peers unaltered.

Module An abstraction used to represent any piece of "code" used to implement a behavior in the JXTA world. Network services are the most common example of behavior that can be instantiated on a peer.

Module Class Represents an expected behavior and an expected binding to support the module; is used primarily to advertise the existence of a behavior.

Module Implementation The implementation of a given module specification; there may be multiple module implementations for a given module specification.

Module Specification Describes a specification of a given module class; it is one approach to providing the functionality that a module class implies. There can be multiple module specifications for a given module class. The module specification is primarily used to access a module.

Peer-to-Peer (P2P)

A decentralized networking paradigm in which distributed nodes, or peers, communicate and work collaboratively to provide services.

Peer A peer is any uniquely identifiable networked device which is capable of interacting with other network devices using standardized protocols.

Peer Group A collection of peers that have a common set of interests and have agreed upon a common set of services.

Peer Resolver Protocol The Peer Resolver Protocol provides a generic query/response interface applications and services can use for building resolution services.

Rendezvous Peer A peer which assumes additional responsibilities for the propagation of messages within a peer group.

Resolver Service Resolver Services are responsible for the routing and processing queries and responses within peer groups. Resolver Services use the Peer Resolver Protocol (PRP) for transmission of queries between peers.

XML (From *Extensible Markup Language (XML) 1.0* "[XML10](#)") : Extensible Markup Language, abbreviated XML, describes a class of data objects called XML documents and partially describes the behavior of computer programs which process them.

Chapter 11

Bibliography

- [ISO11578] *Information Technology - Open Systems Interconnection - Remote Procedure Call (RPC)* , ISO (International Organization for Standardization).
ISO/IEC, 11578:1996.
- [RFC1521] N. Borenstein and N. Freed, *Multipurpose Internet Mail Extensions (MIME) Part One: Mechanisms for Specifying and Describing the Format of Internet Message Bodies* , Internet Engineering Task Force.
IETF RFC, 1521.
- [RFC2045] N. Freed and N. Borenstein, *Multipurpose Internet Mail Extensions (MIME) Part One: Format of Internet Message Bodies* , Internet Engineering Task Force.
IETF RFC, 2045.
- [RFC2046] N. Freed and N. Borenstein, *Multipurpose Internet Mail Extensions (MIME) Part Two: Media Types* , Internet Engineering Task Force.
IETF RFC, 2046.
- [RFC2119] S. Bradner, *Key words for use in RFCs to Indicate Requirement Levels* , Internet Engineering Task Force.
IETF RFC, 2119.
- [RFC2141] R. Moats, *URN Syntax* , Internet Engineering Task Force.
IETF RFC, 2141.
- [RFC2234] P. Overell and D. Crocker, *Augmented BNF for Syntax Specifications: ABNF* , Internet Engineering Task Force.
IETF RFC, 2234.
- [RFC2246] Tim Dierks and Christopher Allen, *The TLS Protocol : Version 1.0* , Internet Engineering Task Force.
IETF RFC, 2246.
- [RFC2616] Roy T. Fielding, James Gettys, Jeffrey C. Mogul, Henrik Frystyk Nielsen, and Larry Masinter, *Hypertext Transfer Protocol -- HTTP/1.1* , Internet Engineering Task Force.
IETF RFC, 2616.
- [RFC3080] Marshall T. Rose, *The Blocks Extensible Exchange Protocol Core* , Internet Engineering Task Force.
IETF RFC, 3080.
- [RFC3081] Marshall T. Rose, *The Blocks Extensible Exchange Protocol Core* , Internet Engineering Task Force.
IETF RFC, 3081.
- [RFC821] Jonathan B. Postel, *Simple Mail Transfer Protocol (SMTP)* , Internet Engineering Task Force.
IETF RFC, 821.

-
- [USA15] Mark Davis and Martin Dürst, *Unicode Standard Annex #15: Unicode Normalization Forms*, The Unicode Consortium.
Unicode Standard Annex, 15.
- [USA28] The Unicode Consortium, *Unicode Standard Annex #28: Unicode 3.2*, The Unicode Consortium.
Unicode Standard Annex, 28.
- [XML10] World Wide Web Consortium, *Extensible Markup Language (XML) 1.0 (Second Edition)*, World Wide Web Consortium.
W3C Recommendation, 2e.
- [XSD2001-1] Henry S. Thompson, David Beech, Murray Maloney, and Noah Mendelsohn, *XML Schema Part 1: Structures*, W3C.
- [XSD2001-2] Paul V. Biron and Ashok Malhotra, *XML Schema Part 2: Datatypes*, W3C.
-